

Information Acquisition and Time - Risk Preference

Daniela O'Donoghue⁰ Weijie Zhong[†]

July 28, 2024

Abstract

An agent acquires information dynamically until her state reaches an upper or lower threshold. She can be subject to a constraint on the rate of entropy reduction by “time risk” — the dispersion of the distribution of her beliefs. We construct a strategy G^* and an ϵ -optimal strategy G^ϵ (minimizing P_{error} and A_{cost}) under either strategy, belief is a compensated Poisson process. In the former, belief that is closer in Bregman divergence. In the latter point with the same entropy as the current belief.

1 Introduction

In this paper, we study information acquisition by a binary state that may be either zero or one. The agent has a belief about the state and is satisfied once her posterior belief reaches either an upper or a lower threshold. She earns a utility from learning and has great flexibility in how she can learn but has a limited budget for her rate of learning. That is, she can choose any probability distribution subject to a constraint on the rate of entropy reduction.

Our simple model captures three important features of learning: flexible learning, limited resources, and thresholds.

⁰odtchen@princeton.edu, Department of Economics, Princeton University

[†]wjzhong@stanford.edu, Graduate School of Business, Stanford University

appear in the contexts of research and development, user-experience testing, and others. For Facebook who must assess whether to introduce a new feature, the unknown state is whether adding the feature increases revenue (and thus profits). The data scientist can learn about the state by running tests. To provide incentives, her manager offers her a bonus that is sufficiently precisely about the state (i.e. her manager offers a bonus that is a minimal level of¹) such that she can sustain a fair expected value design many aspects of the tests—e.g., she can select the subpopulation of users to test the feature and she can adjust the length of time a user has to make other choices. However, there are limits to what she can learn—e.g., her manager does not allow her to run multiple tests simultaneously as doing so could be disastrous for the company, so she can only implement the feature for a given user.

Our main contribution is to show how, in such settings, a strategy that depends on the time that the feature is tested allows for a more accurate representation of preferences over threshold-hitting times beyond the standard discounting. We say that the agent is time-risk neutral if her utility over threshold-hitting times is linear. We show that this is optimal whenever the agent is time-risk neutral, and that it is optimal whenever she is time-risk averse. Critically, these results depend on the shape of the agent's utility function over threshold-hitting times.

In reality, there are many reasons why individuals differ from the predominantly studied case of expected utility. These include external factors such as explicit discounting, opportunity costs associated with foregone opportunities, and internal factors such as present bias resulting from a simple framework that allows for these factors. We derive strategies that are uniformly optimal up to the agent's utility function over threshold-hitting times.

¹In our binary state model, there is a one-to-one mapping between the power of a test (that is, the likelihood of Type I errors) and the bonus.

²In this paper, we model the agent as an expected-utility maximizer. That all of our results will go through as long as the agent's utility over threshold-hitting times is monotonic in the mean-preserving spread order.

We now briefly describe the two learning strategies. When the agent is time-risk averse, Greedy Exploitation is optimal. In this strategy, the agent myopically maximizes the jump to a threshold. She acquires a rare but decisive belief to jump to the threshold that induces her belief to jump to the threshold that is closer. By targeting the closer threshold, she can jump to the threshold without violating her constraint on the rate of entropy reduction. Her belief experiences compensating drift in the direction of the threshold. Eventually, her belief reaches a point that is equal to the two thresholds. At this point, she acquires a belief to jump to either threshold but at rates set so that her belief is stationary in the absence of jumps.

Intuitively, Greedy Exploitation is optimal because of the distribution of threshold-hitting times. Because the probability of an early hitting time is higher, her beliefs drift towards the farther threshold, the jump to the closer threshold is more likely. The amount of time remaining until she hits the closer threshold is smaller than the amount of time remaining until she hits the farther threshold. The agent makes no "progress" in the absence of a jump to the closer threshold. In the absence of late threshold hitting times as well. We, in fact, find that exhaust the agent's resources (in that the constraint is binding at all points in time), Greedy Exploitation achieves the maximal hitting times that is maximal in the mean-preserving sense. Greedy Exploitation maximizes time risk.

When the agent is time-risk averse, she instead chooses Pure Accumulation. In this case, an optimal strategy is to reach a threshold as early as possible. Her beliefs follow a compensating drift process that jumps in the direction of the threshold. If she has an interior belief that has the same entropy as her belief after a jump, her belief experiences a compensating drift in the direction of the threshold. Pure accumulation is a continuous-time analog of the strategy in Ely et al. (2015). The strategy is, in effect, the same as Greedy Exploitation. Because jumps are always to beliefs with the same entropy, the event of a jump, there is no progress: a jump does not change the belief. Instead, all progress is made by jumps.

why the threshold hitting time is deterministic. time risk. It, therefore, produces a distribution mean-preserving spread order among all strategies (in that the constraint on the rate of entropy reduction is time).

Our analysis of optimal learning through the implications for both information acquisition in a model predicts that an agent who is time-risk-loving, whereas an agent who is time-risk-averse is provided the agent has access to these learning strategies agent's space of available learning strategies is generally suboptimal. Thus, when writing models with parameterized signal structures, economists signal structures are without loss of optimality agents they seek to model.

2 Related Literature

Our paper contributes to a large literature on information (1947) and Arrow et al. (1949) we study a sequential the agent to flexibly design the signal process as for (2017), Hébert & Woodford (2023), Steiner et al. (2023). Whereas most of these papers restrict attention to discounting or a linear delay cost, we allow For example, Zhong (2022) assumes exponential discounting preferences whereas Hébert & Woodford (2023) assume a linear delay cost which implies time-risk-loving suggest that the assumed time-risk preferences do not optimal strategies identified in these papers.

³Hébert & Woodford (2023) allow for both discounting and a consider the time-risk neutral limit for the majority of the study how different costs or constraints on information acquisition which is orthogonal to the objective of our paper. Zhong (2022) function that does not have a threshold structure but shows nevertheless similar to Greedy Exploitation.

Pure Accumulation is a continuous-time variant introduced by Ely et al. (2015) in a discrete-time finite time horizon. In Ely et al. (2015), the solution is that it maximizes the expected conditional variance. Harris (2023) also finds that a strategy similar to a setting where the expected payoff is not dependent on the strategy mechanisms behind the optimality of Pure Accumulation. These papers are distinct from that of this paper's aversion.

In our analysis, the key summary statistic that a strategy is the distribution of the time that the agent stops. This statistic is called the *stopping time* and is an object studied in literature on time-risk preferences. Chesson & Viscusi (2003) study the expected discounted utility framework and compare it over time (RSTL) and risk (RATL). DesJardette et al. (2020) show that of models RSTL cannot be violated if there is stopping time experimental evidence suggests that the two are not the same (RATL) (Chesson & Viscusi (2003); Onay & Öncüler (2018) compares both RSTL and RATL and shows that optimal strategies differ dramatically under different time-risk preferences.

The optimal learning strategies that we identify are learning strategies that have been assumed in reduced form literature. For example, Che & Miendorff (2019), Mays & Pans (2018) adopt a framework that restricts attention in order to study optimal stopping with endogenous learning, also often assumed in the literature on strategic learning (Hörner & Skrzypacz (2017)). We show that Poisson learning is a foundation under time ⁵ if the *stopping time* is not exogenous but also related to classic models on the timing of innovation (Stiglitz (1980) and Lee & Wilde (1980) (see a survey by Lee & Wilde (1980) involve a deterministic time of innovation. The

⁴In our paper, the stopping is determined by the chosen learning strategy.

⁵To be clear, we do not show that "all" Poisson learning strategies are optimal strategies in our setting involve Poisson learning.

reduced-form learning process and are non-Bayesian learning strategies in these papers can emerge in an information acquisition framework when agents have time

Our model also allows for Gaussian learning strategies, which are often assumed in reduced-form learning models (Smith (2001); Ke & Villas-Boas (2019); Liang et al. (2021)). Although binary choice problems appear in the literature (Rouder (1998) and Fudenberg et al. (2018)), however, Gaussian learning cannot be justified by optimality except if agents have time-risk neutral preferences provided information

The optimality of a greedy strategy is also the main result. However, the mechanisms in our papers are very different. The optimality crucially depends on the linear-Gaussian setup with multiple information sources and holds for any time preferences and endogenous choice of information sources, but not for exogenous ones. Also related is Gossner et al. (2021) which considers a stopping time (when the belief hits a threshold) in the sense of first-order stochastic dominance.

We model limits on the agent's learning resource by a bound on entropy reduction. That is, the rate of information gain is bounded. This is formally captured by the UPS formulation. The rational inattention literature typically models information costs or constraints by a bound on the rate of information gain (Matijka & McKay (2014); Steiner et al. (2017); Caramazza et al. (2017); Zhong & Bloedel (2021); Morris & Strzalecki (2019)). The information constraint ensures that the expected information gain is bounded for all exhaustive strategies, which is a necessary condition for rational information acquisition. By Theorem 3 in Zhong (2022), the information constraint is both necessary and sufficient for the expected learning to be bounded for all exhaustive strategies.

3 Model

This section presents a simple model of an agent with an unknown state. The unknown state μ_t is a real number in $[0, 1]$. At $t = 0$, the agent believes that $\mu_0 = \mu$ with probability 1. She receives a unit payoff of $\nu^i(t, \mu_t)$ if μ_t is at least μ^i for each $i \in \{1, \dots, N\}$. If μ_t is below μ^i , she receives a payoff of $\lambda^i(t, \mu_t)$. However, she is impatient and her utility is discounted by $\rho(t) = e^{-rt}$. On the other hand, she also experiences a cost $c(\mu_t)$ if μ_t is not equal to μ . The agent's utility is given by

The agent has great flexibility in how she uses her resources and cannot learn about the state μ_t directly. She can only observe the process $\{\mu_t; t \geq 0\}$ that takes values in $[0, 1]$ and satisfies a stochastic differential equation (allowing for jumps) of the form

$$d\mu_t = \sum_{i=1}^N (\nu^i(t, \mu_t) - \mu_t) [dJ_t^i(\lambda^i(t, \mu_t)) - \lambda^i(t, \mu_t) dt] + \sum_{j=1}^M \sigma^j(t, \mu_t) dZ_t^j \quad (1)$$

with $\mu_0 = \mu$ for some positive μ and $\mu^i \in [0, 1]$, and $\{\sigma^j\}_{j=1}^M$ is a standard Brownian motion in \mathbb{R}^M . Above Z_t^j is a standard Brownian motion in \mathbb{R}^M . The process μ_t is a jump-diffusion process. The number of distinct points that the process μ_t visits is the number of distinct Brownian Motions.

We assume that the agent can directly observe μ_t at any time.

$$\mathbb{E} \left[\frac{d}{dt} H(\mu_t) \middle| \mathcal{F}_t \right] \leq I \quad (2)$$

where H is the natural logarithm, I is a constant, and \mathcal{F}_t is the filtration generated by μ_t . Thus, equation (2) is a constraint on the rate of change of $H(\mu_t)$.

⁶Our restriction to jump-diffusion belief processes is within the class of càdlàg processes such that (2) is well defined. This is the case for Harris (2023).

where H is Shannon's entropy so that (2) amounts to a mutual information rate. Without loss of generality, we can assume $\underline{\mu} = 0$ and $\bar{\mu} = 1$. This can be done by redefining

$$\frac{1}{I} \left[H(\mu) - \frac{H(\bar{\mu}) - H(\underline{\mu})}{\bar{\mu} - \underline{\mu}} (\mu - \underline{\mu}) \right].$$

The same belief processes satisfy (2) before and after the transformation, and thus, the drift of the second term is zero. The normalization is convenient because, provided the optional stopping time τ is almost surely finite, the expected time remaining until the process hits the boundary is simply the current entropy:

$$-H(\mu_t) = \mathbb{E}[H(\mu_\tau) - H(\mu_t) | \mathcal{F}_t] = \mathbb{E}[\tau - t | \mathcal{F}_t].$$

To state the agent's belief martingale problem, let τ_μ be a threshold:

$$\tau_\mu := \inf\{t | \mu_t \in [0, \underline{\mu}] \cup [\bar{\mu}, 1]\}.$$

Because μ may be with positive probability for some belief, the agent can stop (or not) at τ_μ and ensure that the agent never hits the boundary.

She solves

$$\max_{\mu \in \mathcal{M}} \mathbb{E}[\rho(\tau_\mu)] \quad (3)$$

such that (2) holds.

Our simple model makes several assumptions in order to distinguish between optimal learning and time-risk preferences in the paper. For example, we assume that the agent expects to live long as (2) is satisfied and that she earns a common return until she ultimately hits. Though these assumptions are not optimal, our model aligns well with a number of economic assumptions. Introduction. A key property of our model that allows it to accommodate general time-risk preferences.

⁷See Theorem 3.22 of Karatzas & Shreve (1998).

⁸That is, we extend the domain and range of

identify optimal learning strategies that apply and to highlight that it is the agent's beliefs that determine the qualitative properties of these strategies. preferences beyond exponential discounting include $\rho(t) = (1 + \alpha t)^{-\gamma/\alpha}$ (Loewenstein & Prelec, 1992); (ii) general utility $u(c) = \phi(e^{-rt})$ where ϕ is increasing and convex (DeJarnett); linear delay cost up to $\rho^T(n) = \max\{T, 0\}$. Time (risk) averse preferences are less commonly assumed in the literature. Experimental evidence that agents may often be time averse (Onay & Öncüler, 2007). As discussed at the end of the paper, time aversion can arise when agents have flow costs of delay.

4 Optimal Learning and Time - Risk

In this section, we present our main results: a strategy that is optimal for a time-risk loving agent and a strategy that is optimal for a time-risk averse agent. These results illustrate the connection between preferences and learning strategies.

4.1 Time - Risk Loving

We first consider the case when the agent is time-risk loving. Her optimal learning strategy is given below in Definition 1. It is informally here.

An optimal strategy for a time-risk loving agent is illustrated in Figure 1.

Let

$$d_H(\tilde{\mu}, \hat{\mu}) = H(\tilde{\mu}) - H(\hat{\mu}) - H'(\hat{\mu})(\tilde{\mu} - \hat{\mu})$$

denote the Bregman divergence between $\tilde{\mu}$ and $\hat{\mu}$ relative to H . In Figure 1, μ^* represents the unique belief that is equal to the divergence to the threshold κ . Initially, the agent either jumps to the threshold that is κ or in Bre

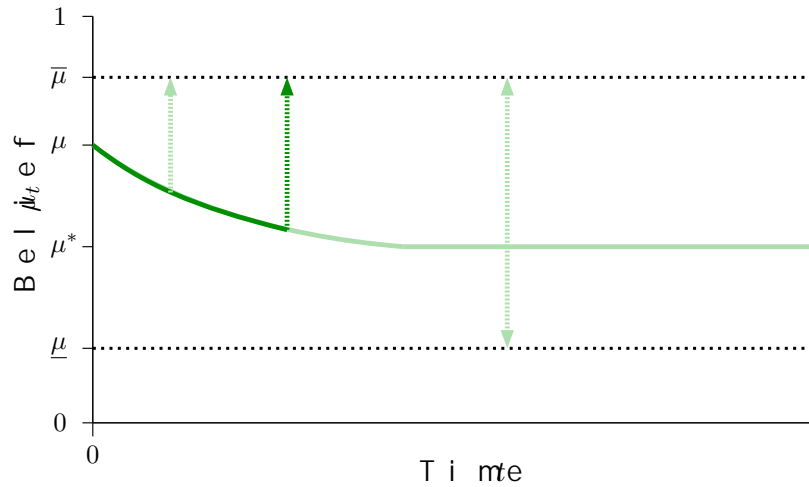


Figure 1: Greedy Exploitation

Notes: In dark green, we plot one possible realization of the Greedy Exploitation strategy. In light green, we plot the belief μ_t^G . The dashed lines with arrows represent jumps in the belief.

experience compensating drift towards the other threshold, the agent greedily maximizes the "change in the "next instant." This is because the agent's belief is μ_t^G when she targets the closer threshold without violating the belief constraint. After some time, in the absence of a jump, the belief μ_t^G converges to a fixed point, her beliefs may jump to either threshold. If the thresholds are such that there is no net compensation for a jump, her beliefs remain stationary.

Definition 6.1 Greedy Exploration strategy. Let $\mu^* \in (0,1)$ be the unique belief $\mu^* \in (0,1)$ such that

- When $\mu_t^G > \mu^*$, her beliefs evolve according to

$$d\mu_t^G = (\bar{\mu} - \mu_t^G) [dJ_t^1(\lambda_t) - \lambda_t dt]$$

where $\lambda_t \in I/d_H(\bar{\mu}, \mu_t^E)$.

- When $\mu_t^G = \mu^*$, her beliefs evolve according to

$$d\mu_t^G = (\bar{\mu} - \mu_t^G) dJ_t^2 \left(\frac{\mu_t^G - \underline{\mu}}{\bar{\mu} - \underline{\mu}} \lambda^* \right) + (\underline{\mu} - \mu_t^G) dJ_t^3 \left(\frac{\bar{\mu} - \mu_t^G}{\bar{\mu} - \underline{\mu}} \lambda^* \right)$$

where $\lambda^* \in 1/d_H(\bar{\mu}, \mu^*)$.

- When $\mu_t^G < \mu^*$, her beliefs evolve according to

$$d\mu_t^G = (\underline{\mu} - \mu_t^G) [dJ_t^1(\lambda_t) - \lambda_t dt]$$

where $\lambda_t \in 1/d_H(\underline{\mu}, \mu_t^G)$.

Above J_t^1 and J_t^2 , and J_t^3 are independent Poisson point processes indicated in parentheses.

Theorem 1. If the agent is time-risk loving > then Kreps

Before we sketch the proof of Theorem 1, we note that the threshold is uniformly optimal for a riskier threshold hitting times among all strategies that at all points in time. To make this precise, we first

Definition 1. $\mathcal{T} = \{ \mu \in \mathcal{M} \text{ such that } (2)_t \text{ binds at all } t \}$

The Greedy Exploitation strategy produces a threshold in the mean-preserving spread order among a set.

Corollary 1. If $\mu \in \mathcal{T}$ then $\tau \in \mathcal{T}$

This result hinges on our assumption that the condition (2). Because of this assumption, all exhaustive threshold hitting time, which is equal to the initial

Proof of Theorem 1 proceeds in seven steps.

Step 1. Set of Basis Functions. We show that any nonnegative ρ can be written as a conical combination of functions

$$\rho_T(t) = \max\{T - t, 0\}$$

where $T \geq 0$. Thus, if Greedy Exploitation is optimal, it must be optimal for any nonnegative dividend process. For any process, it may take on negative values.

Lemma 1 If Greedy Exploitation is optimal for $T \geq 0$, then it is optimal for $T \geq 0$.

See Theorem 3.6 in Müller (1996).

Step 2. Candidate Value Function. Let $V(\mu, T)$ denote the value function for the dividend process X starting at μ at time $T \geq 0$. Let

$$V(\mu, T) = \begin{cases} \int_0^T (T-t) \lambda_t^G e^{-\int_0^t \lambda_z^G dz} dt, & \mu \in (\underline{\mu}, \bar{\mu}) \\ T, & \mu \in \{\underline{\mu}, \bar{\mu}\}. \end{cases} \quad (4)$$

In what follows, it is $\partial V(\mu, T) / \partial \mu = U(\mu, T)$ whenever $\mu \in (\underline{\mu}, \bar{\mu})$ where

$$U(\mu, T) = \int_0^T \lambda_s^G e^{-\int_0^s \lambda_z^G dz} ds$$

is the probability that a dividend is paid by time T starting at μ . To show $\partial V(\mu, T) / \partial \mu \geq 0$ if $\mu \in [\mu^*, \bar{\mu})$ and $\partial V(\mu, T) / \partial \mu \leq 0$ if $\mu \in (\underline{\mu}, \mu^*]$.

To ease the exposition, we adopt $V_T(\mu) = V(\mu, T)$ and $U_T(\mu) = U(\mu, T)$. Also, given a function f , define

$$d_f(\nu, \mu) = f(\nu) - f(\mu) - f'(\mu)(\nu - \mu)$$

where $f'(\mu)$ is well-defined. Note that $d_f(\nu, \mu) \geq 0$ if and only if f is convex at μ .

Step 3. Verify that Greedy Exploitation is optimal. Use the following Lemma 2 which states that the Hamilton-Jacobi-Bellman (HJB) equation (5).

⁹To apply Theorem 3.6 in Müller (1996) easily, it is necessary to show that the dividend process is exhaustive and thus has the expected threshold hitting time.

Lemma 2.10. $V_T(\mu)$ satisfies

$$U_t(\mu) = \max_{\nu} \left\{ \frac{dV_t(\nu, \mu)}{dH(\nu, \mu)}, \frac{V_t''(\mu)}{H''(\mu)} \right\} \quad (5)$$

at $t \in (\underline{\mu}, \bar{\mu}) \times [0, T]$ if $V_T(\mu) \geq 2[m] - \rho^T$.

We first assert that condition (5) is equivalent to

$$U_t(\mu) = \max_{\{\nu^i\}, \{\lambda^i\}, \sigma} \mathcal{A}^{\nu, \lambda, \sigma} V_t(\mu) \quad (6)$$

$$\text{ s.t. } \mathcal{A}^{\nu, \lambda, \sigma} H(\mu_t) \leq 1$$

where $\mathcal{A}^{\nu, \lambda, \sigma}$ is the operator defined by functions

$$\mathcal{A}^{\nu, \lambda, \sigma} f(\mu) = \sum_i \lambda^i d_f(\nu^i, \mu) + \frac{1}{2} \sum_j (\sigma^j)^2 f''(\mu).$$

That $\mathcal{A}^{\nu, \lambda, \sigma}$ is the infinitesimal generator for the process (1). Because it is relatively separable, it suffices to maximize the "bang-for-the-buck" of the drift of H . Therefore, (5) and (6) must be equivalent.

Next, suppose that (5) is satisfied. Consider an $\{\lambda^i\}, \{\sigma^j\}$ with induced first time τ when λ^i is hitting time

$$\begin{aligned} V_T(\mu) &= \mathbb{E} \left[V_{T-\tau \wedge T}(\mu_{\tau \wedge T}) - \int_0^{\tau \wedge T} [-U_{T-t}(\mu_t) + \mathcal{A}^{\nu, \lambda, \sigma} V_{T-s}(\mu_t)] dt \right. \\ &\quad \left. + \sum_j \int_0^{\tau \wedge T} \frac{\partial V_{T-t}(\mu_t)}{\partial \mu} \sigma_t^j dZ_t^j \right. \\ &\quad \left. + \sum_i \int_0^{\tau \wedge T} [V_{T-t}(\nu_t^i) - V_{T-t}(\mu_t)] (dJ_t^i(\lambda_t^i) - \lambda_t^i dt) \right] \\ &= \mathbb{E} \left[V_{T-\tau \wedge T}(\mu_{\tau \wedge T}) - \int_0^{\tau \wedge T} [-U_{T-t}(\mu_t) + \mathcal{A}^{\nu, \lambda, \sigma} V_{T-s}(\mu_t)] dt \right] \\ &\geq \mathbb{E} [V_{T-\tau \wedge T}(\mu_{\tau \wedge T})] \\ &\geq \mathbb{E} [\rho^T(\tau)] \end{aligned}$$

where the first equality uses Itô's formula for j , $\partial V/\partial T = U$ as noted in Step 2, the second equality for $\partial V_{T-t}(\mu_t)/\partial \mu$ and V_{T-t} are bounded which implies that the differentials are true martingales, inequality follows from (6), follows from the definition of \square

Step 4. $\{V_{T-t}(\mu_t^G)\}_{t \in \mathcal{M}}$ — The remaining steps satisfy the conditions of Lemma 2. We begin with Lemma 3 and outer max of (5) is achieved by Greedy Exploitation Lemma. Let $\mu \in [0, \infty)$;

$$\mu \geq \mu^* \quad \mathcal{M}$$

$$U_t(\mu) = \frac{d_{V_t}(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)}.$$

$$\mu \leq \mu^* \quad \mathcal{M}$$

$$U_t(\mu) = \frac{d_{V_t}(\mu, \mu)}{d_H(\mu, \mu)}.$$

Because $V_{T-t}(\mu_t^G) = \mathbb{E}[\rho^T(\tau_{\mu^G}) | \mu_t^G]$ and μ^G is Markov it follows that $\{V_{T-t}(\mu_t^G)\}_{t \in \mathcal{M}}$ is a martingale. For Bayliss's formula, the $\{V_{T-t}(\mu_t^G)\}_{t \in \mathcal{M}}$ is zero if and only if conditions 1 and 2 of the

Step 5. Unimprovable — The proof follows from Lemma 4 showing Greedy Exploitation cannot be improved on by any strategy.

Lemma 4. Let $(\mu, \bar{\mu}) \in (\mu, \bar{\mu}) \times [0, \infty)$;

$$U_t(\mu) = \max_{\nu} \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (7)$$

Proof of Lemma 4. We prove the case $\mu > \bar{\mu}$. The case $\mu < \bar{\mu}$ when μ is analogous. By Lemma 3, it suffices to show that we can split the proof into three cases.

¹ See Theorem 51 of Protter (2005).

- Case $\nu \geq \mu$. We will show that the global derivative of $d_{V_t}(\nu, \mu)/d_H(\nu, \mu)$ in the region $\nu \geq \mu$ starts, we observe that

$$\frac{d}{d\nu} \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)} = \frac{V'_t(\nu) - V'_t(\mu)}{d_H(\nu, \mu)} - \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)^2} [H'(\nu) - H'(\mu)].$$

This derivative is negative if and only if

$$\frac{V'_t(\nu) - V'_t(\mu)}{H'(\nu) - H'(\mu)} \geq \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (8)$$

which is equivalent to

$$\frac{d_{V_t}(\bar{\mu}, \mu) - d_{V_t}(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \geq \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (9)$$

Notice that (9) holds if $d_{V_t}(\bar{\mu}, \mu) \geq d_{V_t}(\bar{\mu}, \nu)$ and $d_H(\bar{\mu}, \mu) \geq d_H(\bar{\mu}, \nu)$. We will show that any local extremum of $d_{V_t}(\nu, \mu)/d_H(\nu, \mu)$ in the region $\nu \geq \mu$ is a local maximum. This immediately implies that $\bar{\mu}$ must be a global maximum.

At any local extremum (9) holds with equality. If $\bar{\mu}$ is a local extremum, then the left-hand side of (9) is negative. This is because the left-hand side is always zero at a local extremum since $d_{V_t}(\bar{\mu}, \mu) = d_{V_t}(\bar{\mu}, \nu)$ and $d_H(\bar{\mu}, \mu) = d_H(\bar{\mu}, \nu)$. The left-hand side of (9) is zero because

$$\begin{aligned} \frac{d}{d\nu} \frac{d_{V_t}(\bar{\mu}, \mu) - d_{V_t}(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} &= \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\nu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &< \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\mu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &= 0 \end{aligned}$$

where we have used $d_{V_t}(\bar{\mu}, \mu) = U_t(\mu)d_H(\bar{\mu}, \mu)$ and $d_{V_t}(\bar{\mu}, \nu) = U_t(\nu)d_H(\bar{\mu}, \nu)$ from Lemma 3 and that $U_t(\nu)$ is increasing in Step 2.

- Case $\nu \leq \mu^*$. In this region, (following the same

easy to show that $d_{V_t}(\bar{\mu}, \mu)/d_H(\bar{\mu}, \mu)$ is nondecreasing if

$$\frac{d_{V_t}(\bar{\mu}, \mu) - d_{V_t}(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \leq \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (10)$$

This is the same condition as (9) except the inequality is reversed. As before, to determine whether a local extremum satisfies the condition, we need to check how the left-hand side is increasing. This can be seen

$$\begin{aligned} \frac{d}{d\nu} \frac{d_{V_t}(\bar{\mu}, \mu) - d_{V_t}(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} &= \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\nu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &> \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\mu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &= 0 \end{aligned}$$

where we have used the fact that the denominator in this region, any local extremum must be a local maximum. $\nu \in (\mu^*, \mu)$ can achieve the maximum in (7).

- Case $\nu \in [\mu, \mu^*]$. Following analogous steps to those above, we find that the derivative $d_{V_t}(\nu, \mu)/d_H(\nu, \mu)$ is positive if and only if

$$\frac{d_{V_t}(\underline{\mu}, \mu) - d_{V_t}(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \mu) - d_H(\underline{\mu}, \nu)} > \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (11)$$

We will prove that the left-hand side of (11) is bounded above by $d_{V_t}(\bar{\mu}, \mu)/d_H(\bar{\mu}, \mu)$. Thus, there can be no $\underline{\mu} \in [\mu, \mu^*]$ that achieves a higher value of $d_{V_t}(\underline{\mu}, \mu)/d_H(\underline{\mu}, \mu)$, since if there was, at $\bar{\mu} = \underline{\mu}$, the derivative would be negative.

To show this, we first observe that

$$d_{V_t}(\underline{\mu}, \mu) = d_{V_t}(\underline{\mu}, \bar{\mu}) + d_{V_t}(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu})(V'_t(\mu) - V'_t(\bar{\mu})), \quad (12)$$

and

$$d_H(\underline{\mu}, \mu) = d_H(\underline{\mu}, \bar{\mu}) + d_H(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu})(H'(\mu) - H'(\bar{\mu})). \quad (13)$$

Define $f(\mu)$ and $g(\mu)$ as

$$f(\mu) = d_{V_t}(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (V'_t(\mu) - V'_t(\bar{\mu})) \quad (14)$$

and

$$g(\mu) = d_H(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (H'(\mu) - H'(\bar{\mu})). \quad (15)$$

Since (8) holds, it follows that

$$\frac{f(\mu)}{g(\mu)} = \frac{d_V(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (V'(\mu) - V'(\bar{\mu}))}{d_H(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (H'(\mu) - H'(\bar{\mu}))} = \frac{d_V(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)}. \quad (16)$$

Also, $d_V(\underline{\mu}, \mu^*)/d_H(\underline{\mu}, \mu^*) =$

$$\frac{f(\mu^*)}{g(\mu^*)} = \frac{d_{V_t}(\underline{\mu}, \bar{\mu}) + f(\mu^*)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu^*)} \Rightarrow \frac{f(\mu^*)}{g(\mu^*)} = \frac{d_{V_t}(\underline{\mu}, \bar{\mu})}{d_H(\underline{\mu}, \bar{\mu})}. \quad (17)$$

Thus,

$$\begin{aligned} \frac{d_{V_t}(\underline{\mu}, \mu) - d_{V_t}(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \mu) - d_H(\underline{\mu}, \nu)} &= \frac{d_{V_t}(\underline{\mu}, \bar{\mu}) + f(\mu) - d_{V_t}(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu) - d_H(\underline{\mu}, \nu)} \\ &= \frac{U_t(\mu^*)d_H(\underline{\mu}, \bar{\mu}) + U_t(\mu)g(\mu) - U_t(\nu)d_H(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu) - d_H(\underline{\mu}, \nu)} \\ &\leq \frac{U_t(\mu^*)d_H(\underline{\mu}, \bar{\mu}) + U_t(\mu)g(\mu) - U_t(\mu^*)d_H(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu) - d_H(\underline{\mu}, \nu)} \\ &\leq U_t(\mu) = \frac{d_{V_t}(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)}. \end{aligned}$$

as desired. The first line uses (12), (13), (14), (16) and (17) and Lemma 3. The $U_t(\nu)$ is added in the numerator for $\nu \in [\underline{\mu}, \mu^*]$ as noted in Step 2.

□

Step 6. Unimprovable by default. Lemma 5 shows Greedy Exploitation cannot be improved on by different

Lemma 5. Let $(\mu, \bar{\mu}) \in (\underline{\mu}, \bar{\mu}) \times [0, \infty)$ and

$$U_t(\mu) \geq \frac{V_t''(\mu)}{H''(\mu)}.$$

Recall from Step 2 that $U_t'(\mu) > 0$ whenever $\mu < \bar{\mu}$. Then

$$U_t'(\mu) = \frac{d}{d\mu} \frac{dV_t(\bar{\mu}, \mu)}{dH(\bar{\mu}, \mu)} = \frac{-d_H(\bar{\mu}, \mu)V_t''(\mu)(\bar{\mu} - \mu) + d_{V_t}(\bar{\mu}, \mu)H''(\mu)(\bar{\mu} - \mu)}{d_H(\bar{\mu}, \mu)^2} > 0$$

which implies that

$$U_t(\mu) = \frac{dV_t(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)} > \frac{V_t''(\mu)}{H''(\mu)}$$

as desired. An analogous argument works when $\mu = \bar{\mu}$ follows from continuity. \square

Step 7. Putting Lemma 4 together imply that (5) is true. Lemma 2 then implies that Greedy Exploitation discount function of Lemma 1 then implies optimal convex. The proof of Theorem 1 is complete. \square

4.2 Time - Risk Averse

When the agent is time-risk averse, the Accumulation illustrated graphically below in Figure 2.

As discussed in Section 2, the Pure Accumulation is a suspense-maximizing strategy (2015). Under this belief either jumps in the direction of the farthest satiating drift. When her belief jumps, it jumps to her current belief so that all progress is made through

Definition 3. Pure Accumulation strategy is defined on $[0, 1] \setminus \{\mu^*\}$ as follows. Let $\hat{\mu}: [0, 1] \rightarrow [0, 1]$ denote the function on the interval $[0, 1]$ such that $H(\hat{\mu}(\mu)) = H(\mu)$. Under Pure accumulation, the agent




Figure 2: Pure Accumulation

Note the dark green curve represents μ_t^P on the possible belief segments represent jumps. The light green curves represent path μ_t^H of

according to

$$d\mu_t^P = [\mu^H(\mu_t^P) - \mu_t^P] dJ_t(\lambda_t) - \lambda_t [\mu^H(\mu_t^P) - \mu_t^P] dt$$

where J_t is a Poisson point process with $\lambda_t = \lambda_H(\mu_t^H, \mu_t^P)$ ticks at rate λ_t . Theorem 2: If the agent is time-risk averse > then Pure Accumulation. Proof: Under Pure Accumulation, the agent is guaranteed to determine $\mu^H(\mu)$ time $\tau \in \mathcal{T}$. \square

Because Pure Accumulation entails no time risk, the result follows.

Corollary 2.01 states that $\tau \in \mathcal{T}$

5 Concluding Discussion

In this paper, we have studied the relationship between time and information acquisition. We have shown that

loving agent is Greedy Exploitation. This strategy over threshold hitting times among all exhaustive optimal strategy for a time-risk averse agent is produces a deterministic threshold hitting time of these strategies are uniformly optimal up to the utility function, provided the agent is impatient. This inconsistency. In practice, agents may have time well-studied case of exponential discounting. How these agents may seek to acquire information economists may consider using when modeling these

In order to illustrate the connection between learning as sharply as possible we have made a number of specifications of binary states, fixed stopping thresholds, speed are critical because they ensure that all expected threshold hitting times are the same. Why Greedy Exploitation and Pure Accumulation are optimal is because they remain minimal threshold-hitting times in the mean-pre-haustive strategies. This allows us to emphasize determines their optimality. The assumption that threshold hitting times and not on which threshold is hit at all times but is not critical for the economic insight into the qualitative properties of Greedy Exploitation under other model formulations, for example, with states, and endogenously chosen stopping thresholds. Examples in the literature where this is so as reviewed possible to extend our model of optimal learning to environments though explicit.¹² The advantage of our special setup is that it is possible to solve for optimal uniformly optimal for large classes of payoffs and

¹Specifically, with a binary state and fixed thresholds, even probability distribution over terminal beliefs by the market with multiple states. That is, all learning strategies yield the same constraint then ensures that all learning strategies that yield the same expected threshold-hitting times.

²Our solutions were based on a guess and verify approach that exploits the structure of our setup. In more general setups guessing the

role of time risk for optimal learning.

There are two promising avenues to explore in future work: how our results may extend to the case when the agent is time-risk averse. For these more general preferences, what is the optimal information acquisition? A second avenue is to extend our model of information acquisition into strategic settings with multiple agents in order to study the implications of flexible information acquisition.

References

- Arrow, K. J., Blackwell, D., & Girshick, M. A. (1953). Bayes' theorem and the theory of sequential decisions. *Journal of the American Statistical Association*, 48(262), 213–244.
- Caplin, A., Dean, M., & Leahy, J. (2017). Time discounting and time preference: characterizing and generalizing. *Working Paper 2365*, National Bureau of Economic Research.
- Che, Y.-K., & Miendorff, K. (2019). Optimal dynamic information acquisition. *American Economic Review*, 109(18), 2993–3029.
- Chen, M. K. (2013). The effect of language on economic savings rates, health behaviors, and economic growth. *American Economic Review*, 103(2), 690–731.
- Chesson, H. W., & Viscusi, W. K. (2003). Commonalities in risk and time discounting. *Journal of Human Capital*, 5(1), 57–71.
- Dasgupta, P., & Stiglitz, J. (1980). Uncertainty, industrial structure, and the nature of innovative activity. *The Bell Journal of Economics*, 1(2), 1–28.
- DeJarnette, P., Dillenberger, D., Gottlieb, D., & Smith, J. (2018). Time discounting and time preference: characterizing and generalizing. *Working Paper 2365*, National Bureau of Economic Research.
- Ely, J., Frankel, A., & Kamenica, E. (2015). The role of time risk for optimal learning. *Journal of Political Economy*, 123(1), 215–260.

- Frankel, A. & Kamenica, E. (2019). Quasi-rationality in the American Economic Review 3650–80.
- Fudenberg, D., Strack, P., & Strzalecki, T. (2018). Timing of American Economic Review 121, 651–84.
- Georgiadis-Harris, A. (2023). Information acquisition and rationality. *Journal of Economic Theory* 201, 10244.
- Gossner, O., Steiner, J., & Stewart, M. (2012). *Journal of Economic Theory* 146(4), 1717–1751.
- Hébert, B. & Woodford, M. (2017). Attention and sequential sampling. Technical report, National Bureau of Economic Research.
- Hébert, B. & Woodford, M. (2023). Rational inattention. *Journal of Economic Theory* 201, 10244.
- Hörner, J. & Skrzypacz, A. (2017). Learning, experimentation, and rationality. *Advances in Econometrics and Econometrics* 31, 1–30.
- Karatzas, I. & Shreve, S. E. (1998). *Brownian motion and stochastic calculus* (pp. 47–127).
- Ke, T. T. & Villas-Boas, J. M. (2019). *Journal of Economic Theory* 188, 383–437.
- Lee, T. & Wilde, L. L. (1980). Market structure and the quarterly journal of economics 94(2), 429–436.
- Liang, A., Mu, X., & Syrgkanis, V. (2019). Dynamic rationality. *arXiv preprint arXiv:1910.07015*.
- Loewenstein, G. & Prelec, D. (1992). Anomalies in time discounting and time preference. *Quarterly Journal of Economics* 107(2), 557–593.
- Matijka, F. & McKay, A. (2014). Rational inattention: a foundation for the multi-armed bandit problem. *Journal of Economic Theory* 146(1), 272–298.
- Mayskaya, T. (2019). Dynamic choice of information. manuscript.

- Morris, S. & Strack, W. (2001). The Equivalence of Sampling and Static Weibull Risk Premium Costs. *Economic Journal*, 111(6), 1629–1644.
- Moscarini, G. & Smith, L. (2001). The Equivalence of Sampling and Static Weibull Risk Premium Costs. *Economic Journal*, 111(6), 1629–1644.
- Müller, A. (1996). Orderings of risks: A comparative insurance: Mathematical Economics. *Mathematical Economics*, 17(3), 215–222.
- Nikandrova, A. & Pans, R. (2018). The Dynamic Expected Utility. *Economic Journal*, 118(1), 115–143.
- Onay, S. & Öncüler, A. (2007). Intertemporal choice and the role of risk aversion. *Journal of Risk and Uncertainty*, 24(1), 99–112.
- Protter, P. (2005). Stochastic differential equations. *Stochastic Processes and their Applications*, 115(1), 1–14.
- Ratcliff, R. & Rouder, J. N. (1998). Modeling response times in the diffusion model. *Psychological Review*, 105(1), 34–81.
- Reinganum, J. F. (1989). The timing of innovation. *Handbook of industrial organization*, 2, 641–706.
- Sims, C. A. (2003). Implications of a rational expectations model for macroeconomics. *Economic Journal*, 113(3), 665–690.
- Steiner, J., Stewart, C., & Matijka, F. (2017). Risk and delay in decision making. *Journal of Economic Psychology*, 65, 521–553.
- Wald, A. (1947). Foundations of a general theory of econometrics. *Journal of the American Statistical Association*, 42(238), 208–237.
- Zhong, W. (2022). Optimal dynamic pricing in a market with learning. *Journal of Economic Theory*, 202, 11537–11582.
- Zhong, W. & Bloedel, A. (2021). The cost of optimal pricing. *Journal of Economic Theory*, 202, 11537–11582.