

Information Acquisition and Time-Risk

By Daniel Chen and Wei Jie Zhong

An agent acquires information dynamically until a binary state reaches an upper or lower threshold of any signal process subject to a constraint on total resource reduction. Strategies are ordered by "time risk" based on the distribution of threshold-hitting times. We compare maximizing time risk (Greedy Exploitation) and minimizing it (Pure Accumulation). Under either strategy, the process is a compensated Poisson process. In the former, the threshold is closer in Bregman divergence. In the latter, it is closer to the unique point with the same entropy as the JEL: D80, D81, D83

Keywords: Information acquisition, stopping time, Bregman divergence, rational inattention

Introduction

In this paper, we study information acquisition by an agent with an unknown binary state that may be either zero or one. The agent is reasonably certain about the state and is satisfied with the state if it reaches either an upper or a lower threshold. She has a great flexibility in learning, but a resource constraint that limits her rate of learning. Her posterior belief process is subject to a constraint on total resource reduction. Our simple model captures three important features: flexible learning, limited resources, and thresholds. These often appear in the contexts of research and development, marketing, user-experience testing, and others. For example, a data scientist at Facebook who must assess whether to introduce a new feature on the platform. The unknown state is whether adding the feature will increase user engagement (and thus profits). The data scientist learns about the state by conducting A/B tests. To provide incentives, her compensation is based on the state when she has learned sufficiently precisely about the state. She must identify the state at some minimum time. She can design many aspects of the A/B tests—e.g., she can

Chen: dtchen@princeton.edu, Department of Economics, Princeton University
jie.zhong@stanford.edu, Graduate School of Business, Stanford University
We thank excellent research assistance.

¹In our binary state model, there is a one-to-one mapping between the power of a test (that is, the likelihood of Type II errors).

users that she gives access to the feature and she can user has access to the feature among other choices. what she can do that constrain her speed of learning allow her to implement the feature for all users simultaneously be disastrous for profits if the feature is disliked feature for a given user for a short amount of time, e

Our main contribution is to show how, in such settings learning strategy time-period risk on the feature allows for a rich set of preferences over threshold-hitting times of exponential discounting. We say that the agent is whenever her utility over threshold-hitting times is a learning strategy that is optimal whenever the agent strategy that is optimal whenever she is time-risk aware of these strategies does not depend on the shape of beyond its convexity or concavity.

In reality, there are many reasons why individuals that differ from the predominantly studied case of exponential discounting may be due to external factors such as explicit deadlines, and flow costs associated with foregone opportunities may also be due to internal factors such as present bias and discounting. We provide a simple framework that allows studying optimal learning and derive strategies that depend on the convexity or concavity of the utility function on discounting.

We now briefly describe the two learning strategies that we study. When the agent is Greedy Exploitation is optimal. Under this strategy, the agent myopically chooses the rate that her beliefs jump to a threshold. She acquires information that, upon arrival, induces her belief to jump to the threshold. the Bregman divergence (under the entropy function on the threshold, she can jump at a faster rate without violating the rate of entropy reduction. In the absence of a signal compensating drift in the direction of the farther threshold, belief reaches a point that is equidistant in the Bregman divergence from the thresholds. At this point, she acquires signals such that either threshold but at rates set so that there is no drift. her belief is stationary in the absence of a jump.

Intuitively, Greedy Exploitation is optimal because of the distribution of threshold-hitting times. Because of a high probability of an early hitting time. However, since beliefs drift towards the farther threshold,

²In this paper, we model the agent as an expected-utility maximizer. All of our results will go through as long as the agent has a preference that is monotonic in the mean-preserving spread order.

the expected amount of time remaining until a threshold is reached. In this sense, the agent makes no "progress" in the sense that there is a high probability of late threshold hitting. We show that among all strategies that exhaust the agent's time budget, Greedy Exploitation yields a distribution of hitting times that is mean-preserving spread order. In this sense, it maximizes the expected amount of time remaining until a threshold is reached.

When the agent is time-risk averse, she instead se-
 In this case, an optimal cumulative strategy, he
 beliefs reach a threshold that her beliefs follow a comp-
 Poisson process that jumps in the direction of the t-
 but to an interior belief that has the same entropy
 the absence of a jump, her belief experiences a com-
 closer threshold. Pure accumulation is a continuous
 maximal " policy in Ely, Frankel and Kamenica (2015).
 the " opposite " of Greedy Exploitation. Because jum-
 the same entropy as the current belief, in the event of
 a jump does not reduce the expected amount of time u-
 Instead, all progress is made through drift, which
 time is deterministic. Thus at Pure Accumulation, the mean drift
 produces a distribution of hitting times that is mi-
 spread order among all strategies that exhaust the a-
 constraint on the rate of entropy reduction is bindi-

Our analysis of optimal learning through the lens of implications for both information acquisition in p. Our model predicts that an agent who is time-risk-lo exploitation, whereas an agent who is time-risk-avers lation, provided the agent has access to these learn ever the agent's space of available learning strateg signals are generally suboptimal. Thus, when writi information with parameterized signal structures, whether these signal structures are without loss of preferences of the agents they seek to model.

1. Related Literature

Our paper contributes to a large literature on inf Wald (1947) and Arrow, Blackwell and Girshick (194 sampling problem but allow the agent to flexibly des in Zhong (2022), Hébert and Woodford (2017), Héber Steiner, Stewart and Matjka (2017), and Georgiad most of these papers restrict attention to the stan counting or a linear delay cost, we allow for more ge example, Zhong (2022) assumes exponential discount

loving preferences whereas Hébert and Woodford (2023) assume a linear delay cost which implies. Our results suggest that the assumed time-risk preferences of the optimal strategies identified in the

Pure Accumulation is a continuous-time variant of the strategy introduced by Ely, Frankel and Kamenica (2015) with a deterministic and finite time horizon. In Ely, the suspense-maximal strategy is optimal in that it is that is increasing in the variances of future beliefs. It finds that a strategy similar to Pure Accumulation is the stopping strategy does not depend on the learning. The mechanisms behind the optimality of Pure Accumulation papers are distinct from that of this paper where it is aversion.

In our analysis, the key summary statistic that determines a strategy is the distribution of the time that the strategy crosses a threshold. This statistic is the focus of an object study in an emerging literature on time-risk preferences and Chen (2013) show that the expected discounted utility preferences is increasing over (R, S, T, θ) . In Detemmer et al. (2020) show that within a broad class of models R, S, T, θ there is stochastic impatience. However, experimental subjects raise risk aversion over (R, S, T, θ) (Chen and Viscusi (2003); Onay and Öncüler (2007)). Our model accommodates RATL and shows that optimal information acquisition under different time-risk preferences.

The optimal learning strategies that we identify are learning strategies that have been assumed in reduced form in the literature. For example, Che and Miendorff and Nikandrova and Pans (2018) adopt a framework that uses Poisson signal processes in order to study optimal information formation. Poisson signals are also often assumed in experimentation (see a survey by Hörner and Skrzypacz). Poisson learning has an optimization foundation under uncertainty⁵. The Pure Accumulation strategy is also related

³Hébert and Woodford (2023) allow for both discounting and a linear cost. They consider the time-risk neutral limit for the majority of their analysis. However, different costs or constraints on information acquisition affect optimal strategies orthogonal to the objective of our paper. Zhong (2022) assumes a belief that does not have a threshold structure but shows that the optimal learning strategy is Greedy Exploitation.

⁴In our paper, the stopping is determined by the chosen learning strategy.

⁵To be clear, we do not show that the learning strategies are optimal but rather that the strategies in our setting involve Poisson learning.

timing of innovation introduced by Dasgupta and Stiglitz (1980) (see a survey by Reinganum (1989)) which determines the time of innovation. The models in these papers assume a sequential learning process and are non-Bayesian. However, we show that the models in these papers can emerge endogenously in a Bayesian framework when agents have time risk-averse preferences.

Our model also allows for Gaussian learning strategies. Gaussian processes are often assumed in reduced-form learning models (e.g., Morris and Strack (2019); Ke and Villas-Boas (2019); Liang, Morris and Strack (2019)). In a Bayesian model, Gaussian learning problems appear in Ratcliff and Rouder (1998) and Fudenberg and Leckie (2018). However, our results imply that Gaussian learning is optimal only in the knife-edge case when preferences provided information can be acquired for free.

The optimality of a greedy strategy is also the main result of Syrgkanis (2019). However, the mechanisms in our paper are different. Liang, Mu and Syrgkanis (2019)'s result crucially depends on the setup with exogenously given Gaussian information sources and preferences. Our result allows for a flexible and general learning setup with multiple information sources, but crucially depends on time preferences. Steiner and Stewart (2021) which studies a model with multiple information sources and derives a greedy learning strategy that achieves the shortest stopping time (when the belief hits an exogenous threshold) in the sense of first-order stochastic dominance.

We model limits on the agent's learning resources via a constraint on entropy reduction. That is, the rate of resource consumption is uniformly bounded. This is a natural constraint in the rational inattention literature (e.g., Matějka and McKay (2014); Steiner, Stewart, Dean and Leahy (2017)). Microfoundations for the UPS in Frankel and Kamenica (2019); Caplin, Dean and Leahy (2021); Bloedel (2021); Morris and Strack (2019). In our paper, the constraint ensures that the expected threshold hit by an exhaustive strategy, which is the expected time to acquire information, is finite. By Theorem 3 in Zhong (2022), a UPS is both necessary and sufficient for the expected learning time to be finite for all exhaustive strategies.

11. Model

This section presents a simple model of an agent who learns about an unknown state $\omega \in \Omega$. The unknown state ω is drawn from a distribution $\mu \in \Delta(\Omega)$. The agent believes ω is drawn from μ . She receives a unit payoff when her posterior belief μ_t is either an upper threshold $\bar{\mu}$ or a lower threshold $\underline{\mu}$.

$\bar{\mu} \in (\mu, 1)$ or a lower threshold. However, she is impatient and utility is a decreasing function of the hitting time ρ is convex, we say that the risk aversion is concave we say that the risk aversion is convex arises when the agent discounts time at rate ρ . On the other hand, a concave utility is if the agent's utility increases as the discount rate increases.

The agent has great flexibility but has a limited number of resources and cannot learn from the past. Let $\mu = \{\mu_t; t \geq 0\}$ that take values and satisfy a stochastic differential equation (allowing for jumps) of the form

$$(1) \quad d\mu_t = \sum_{i=1}^N (\nu^i(t, \mu_t) - \mu_t) [dJ_t^i(\lambda^i(t, \mu_t)) - \lambda^i(t, \mu_t) dt] + \sum_{j=1}^M \sigma^j(t, \mu_t) dZ_t^j$$

with $\mu_0 = \mu$ for some positive initial condition μ and $\lambda^i(t, \mu_t) \geq 0$ and $\sigma^j(t, \mu_t) \geq 0$. Above Z_t^j is a standard Brownian motion and J_t^i is a Poisson point process with intensity $\lambda^i(t, \mu_t)$. If J_t^i has n jumps to μ_t then the number of distinct points that J_t^i jumps to at time t is n . The number of distinct Brownian motions is the integer M .

We assume that the agent can directly choose any belief that for all

$$(2) \quad \mathcal{A}H(\mu_t) := \lim_{s \rightarrow 0} \mathbb{E} \left[\frac{H(\mu_{t+s}) - H(\mu_t)}{s} \middle| \mathcal{F}_t \right] \leq I,$$

where $\{\mathcal{F}_t\}$ is the natural filtration of μ_t and H is a strictly convex function defined on $(0, 1)$ and $I \geq 0$ is a constant. Thus, the agent's belief to associate with μ_t is $H(\mu_t)$. Equation (2) is a constraint on the reduction. A specific example is when $H(\mu) = -\log(\mu)$ so that (2) is a constraint on the well-known mutual information. In general, $H(\mu) = -\log(\mu)$ is a simplification. This can be done by redefining

$$\frac{1}{I} \left[H(\mu) - \frac{H(\bar{\mu}) - H(\underline{\mu})}{\bar{\mu} - \underline{\mu}} (\mu - \underline{\mu}) \right].$$

The same belief processes satisfy (2) before and after the beliefs are martingales (and thus, the drift of the

⁶Our restriction to jump-diffusion belief processes is without loss of generality. We can restrict to jump-diffusion belief processes such that (2) is well defined. This follows from Proposition 2.3 of (2023).

is zero).

To state the agent's learning problem, let the agent's beliefs reach a threshold:

$$\tau_\mu := \inf\{t | \mu_t \in [0, \underline{\mu}] \cup [\bar{\mu}, 1]\}.$$

We extend the domain ρ such that $\rho(\infty) = 0$ to ensure that the agent never stops learning (i.e., τ_μ may be ∞ if beliefs do not reach a positive probability). The agent solves

$$(3) \quad \max_{\mu \in \mathcal{M}} \mathbb{E}[\rho(\tau_\mu)]$$

such that (2) holds.

Our simple model makes several assumptions in order to bridge between optimal learning and time-risk preferences in this paper. For example, we assume that the agent experiences a continuous stream of information (as long as (2) is satisfied) and that she earns a continuous reward (as long as (2) is satisfied). Though these assumptions may seem restrictive, we believe that our model aligns well with a number of empirical findings discussed in the Introduction. A key property of our model is that it accommodates general time-risk preferences. Our contribution is to identify optimal learning strategies for time-risk loving preferences and to highlight the role of time-risk preferences in determining the qualitative properties of these strategies. We consider time-risk loving preferences beyond exponential discounting (i.e., $\rho(t) = \gamma^t$) (e.g., Leowenstien and Prelec, 1992) and generalized expected utility (e.g., DeJarnette et al., 2020); and linear delay costs (i.e., $\rho^T(t) = \max\{T - t, 0\}$). Time-risk averse preferences are less common in the literature though there is growing experimental evidence that people are often time-risk averse (Chesson and Viscusi, 2002). As discussed at the start of this section, time-risk loving preferences have low costs of delay that increase with time.

III. Optimal Learning and Time-Risk Preferences

In this section, we present our main results: a strategy that is optimal for the agent if she is time-risk loving and a strategy that is optimal if she is time-risk averse. These results illustrate the connection between optimal learning and time-risk preferences.

A. Time-Risk Loving

We first consider the case when the agent is time-risk loving. In this case, her optimal learning strategy is given below in D.

describing it informally here.

An optimal strategy for greedy exploitation is illustrated in Figure 1.




Figure 1. Greedy Exploitation

Notes: In dark green, we plot one possible greedy belief path. In light green, we plot the threshold belief path. Dashed lines with arrows represent jumps in the belief path. The figure is computed using $\mu = 0.5$, $\mu^* = 0.3$, and $\bar{\mu} = 0.7$.

Let

$$d_H(\tilde{\mu}, \hat{\mu}) = H(\tilde{\mu}) - H(\hat{\mu}) - H'(\hat{\mu})(\tilde{\mu} - \hat{\mu})$$

denote the Bregman divergence between two beliefs $\tilde{\mu}$ and $\hat{\mu}$. In Figure 1, μ^* represents the unique belief that is equal to the Bregman divergence to μ and μ^* is the unique belief that is equal to the Bregman divergence to μ . In this case, as experience compensating drift towards μ^* is targeted, the agent greedily updates her beliefs reach a threshold in the "next instant". The belief jumps can jump at a faster rate when she targets μ^* by violating her resource constraint (2). After some time, her beliefs eventually reach μ^* . At this point, her beliefs may jump to μ threshold. The jump rates to the respective thresholds are equal, and so, in the absence of compensating drift, the belief is stationary.

DEFINITION 4: Greedy Exploitation is defined as follows. Let $\mu^* \in (0, 1)$ be the unique belief μ^* such that $d_H(\mu, \mu^*) = d_H(\mu, \mu^*)$, that

- When $\mu_t^G > \mu^*$, her beliefs evolve according to

$$d\mu_t^G = (\bar{\mu} - \mu_t^G) [dJ_t^1(\lambda_t) - \lambda_t dt]$$

where $\lambda_t \in I/d_H(\bar{\mu}, \mu_t^G)$.

- When $\mu_t^G = \mu^*$, her beliefs evolve according to

$$d\mu_t^G = (\bar{\mu} - \mu_t^G) dJ_t^2 \left(\frac{\mu_t^G - \underline{\mu}}{\bar{\mu} - \underline{\mu}} \lambda^* \right) + (\underline{\mu} - \mu_t^G) dJ_t^3 \left(\frac{\bar{\mu} - \mu_t^G}{\bar{\mu} - \underline{\mu}} \lambda^* \right)$$

where $\lambda^* \in 1/d_H(\bar{\mu}, \mu^*)$.

- When $\mu_t^G < \mu^*$, her beliefs evolve according to

$$d\mu_t^G = (\underline{\mu} - \mu_t^G) [dJ_t^1(\lambda_t) - \lambda_t dt]$$

where $\lambda_t \in 1/d_H(\underline{\mu}, \mu_t^G)$.

Above J_t^1 , and J_t^2 are independent Poisson point processes indicated in parentheses.

THEOREM 1 If the agent is time-risk loving, then Greedy is optimal.

Before we sketch the proof of Theorem 1, we note that this result is uniformly optimal in the sense that it is the best strategy of threshold hitting times among all strategies that are satisfied at any point before stopping. To make this precise, we give the following definition.

DEFINITION 1 A strategy τ_μ is μ -optimal if for all $\mu \in \mathcal{M}$, $\mathbb{E}[H(\mu_{\tau_\mu})] \geq \mathbb{E}[H(\mu_t)]$ for all $t \geq 0$.

The following Lemma 1 implies that the expected value of the entropy is equal to the initial entropy for any exhaustive strategy.

LEMMA 1 For $\mu \in \mathcal{M}$, $\mathbb{E}[H(\mu_{\tau_\mu})] \geq H(\mu)$ with equality if and only if τ_μ is μ -optimal.

PROOF:

Recall that we have $H(\mu_t) = H(\mu) + \int_0^t \lambda_s ds$. Therefore,

$$\begin{aligned} -H(\mu) &= \mathbb{E}[H(\mu_{\tau_\mu}) - H(\mu)] \\ &= \mathbb{E} \left[\int_0^{\tau_\mu} \lambda_t dt \right] \\ &\leq \mathbb{E} \left[\int_0^{\infty} \lambda_t dt \right] \\ &= \mathbb{E}[\tau_\mu] \end{aligned}$$

where the second equality follows from Dynkin's formula (2005) which has a martingale and the optional stopping theorem (see Theorem 3.22 of Karatzas and Shreve (1998)). The inequality follows from the constraint (2). The inequality is strict if $\tau \in \mathcal{T}$.

Theorem 1 and Lemma 1 together imply that the Greedy strategy produces a threshold-hitting time that is maximal in order among all threshold-hitting times in this set.

COROLLARY 1 *Let $\tau \in \mathcal{T}$ and $\tau \leq \tau^*$. Then*

PROOF OF THEOREM 1:

The proof proceeds in seven steps.

Step 1. Set of Basis Functions. Let $\rho_T(t)$ be a nonnegative convex function on $[0, T]$ that can be written as a conical combination of functions

$$\rho_T(t) = \max\{T - t, 0\}$$

where $T \geq 0$. Thus, if Greedy is optimal for ρ_T , it must be optimal for any nonnegative convex function ρ that can be written as a conical combination of functions of the form ρ_T .

LEMMA 2: If Greedy is optimal for ρ_T for all $T \geq 0$, then it is optimal for any convex function ρ .

PROOF OF LEMMA 2:

See Theorem 3.6 in Müller (1996).

Step 2. Candidate Value Function. Let $V(\mu, T)$ denote the value function for the Greedy strategy for ρ_T and $U(\mu, T)$ denote the value function for the optimal strategy for ρ_T . For $T \geq 0$, let

$$(4) \quad V(\mu, T) = \begin{cases} \int_0^T (T-t) \lambda_t^G e^{-\int_0^t \lambda_z^G dz} dt, & \mu \in (\underline{\mu}, \bar{\mu}) \\ T, & \mu \in \{\underline{\mu}, \bar{\mu}\}. \end{cases}$$

In what follows, it is assumed that $\partial V(\mu, T) / \partial T = U(\mu, T)$ where $\mu \in (\underline{\mu}, \bar{\mu})$ where

$$U(\mu, T) = \int_0^T \lambda_t^G e^{-\int_0^t \lambda_z^G dz} dt$$

is the probability that T is hit by Greedy before $\bar{\mu}$ is hit. It is easy to see that $\partial V(\mu, T) / \partial \mu > 0$ if $\mu \in [\mu^*, \bar{\mu})$ and $\partial V(\mu, T) / \partial \mu < 0$ if $\mu \in (\underline{\mu}, \mu^*]$.

⁷To apply Theorem 3.6 in Müller (1996) we need to show that the Greedy strategy is exhaustive and thus has the expected threshold-hitting time of

To ease the exposition, we adopt the following notation $V(\mu, T)$ and $d_T(\mu) = U(\mu, T)$. Also, given any μ and ν , we let

$$d_f(\nu, \mu) = f(\nu) - f(\mu) - f'(\mu)(\nu - \mu)$$

when $f(\mu)$ is well defined. Note that $d_f(\nu, \mu)$ is the Bregman divergence.

Step 3. Verify that the optimality of Greedy E we use the following Lemma 3 which states that the Hamilton - Jacobi - Bellman (HJB) equation (5).

LEMMA 3 Given $\rho \geq 0$, if U_t satisfies

$$(5) \quad U_t(\mu) = \max \left\{ \max_{\nu} \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}, \frac{V_t''(\mu)}{H''(\mu)} \right\}$$

at $(\mu, t) \in (\underline{\mu}, \bar{\mu}) \times [0, T]$ then $U_t(\mu)$ is equal to ρ^T .

PROOF OF LEMMA 3:

We first assert that condition (5) is equivalent to

$$(6) \quad U_t(\mu) = \max_{\{\nu^i\}, \{\lambda^i\}, \sigma} \mathcal{A}^{\nu, \lambda, \sigma} V_t(\mu) \\ \text{s.t. } \mathcal{A}^{\nu, \lambda, \sigma} H(\mu_t) \leq 1$$

where $\mathcal{A}^{\nu, \lambda, \sigma}$ is the operator defined on $C^2(\underline{\mu}, \bar{\mu})$ by functions

$$\mathcal{A}^{\nu, \lambda, \sigma} f(\mu) = \sum_i \lambda^i d_f(\nu^i, \mu) + \frac{1}{2} \sum_j (\sigma^j)^2 f''(\mu).$$

That $\mathcal{A}^{\nu, \lambda, \sigma}$ is the infinitesimal generator for the controlled process (1). Because $\mathcal{A}^{\nu, \lambda, \sigma}$ is additively separable, it succeeds to jump point or volatility to achieve the max in (6). σ is chosen to maximize the "bang-for-the-buck" — that is, V to the d_H . Therefore, (5) and (6) must be equivalent.

Next, suppose that (5) is satisfied. Consider an ar

$\{\nu^i\}, \{\lambda^i\}, \{\sigma^j\}$ with induced first three Webladve hitting time

$$\begin{aligned}
 V_T(\mu) &= \mathbb{E} \left[V_{T-\tau \wedge T}(\mu_{\tau \wedge T}) - \int_0^{\tau \wedge T} \left[\frac{\partial V_{T-t}}{\partial t}(\mu_t) + \mathcal{A}^{\nu, \lambda, \sigma} V_{T-t}(\mu_t) \right] dt \right. \\
 &\quad + \sum_j \int_0^{\tau \wedge T} \frac{\partial V_{T-t}(\mu_t)}{\partial \mu} \sigma_t^j dZ_t^j \\
 &\quad \left. + \sum_i \int_0^{\tau \wedge T} [V_{T-t}(\nu_t^i) - V_{T-t}(\mu_t)] (dJ_t^i(\lambda_t^i) - \lambda_t^i dt) \right] \\
 &= \mathbb{E} \left[V_{T-\tau \wedge T}(\mu_{\tau \wedge T}) - \int_0^{\tau \wedge T} \left[\frac{\partial V_{T-t}}{\partial t}(\mu_t) + \mathcal{A}^{\nu, \lambda, \sigma} V_{T-t}(\mu_t) \right] dt \right] \\
 &= \mathbb{E} \left[V_{T-\tau \wedge T}(\mu_{\tau \wedge T}) - \int_0^{\tau \wedge T} [-U_{T-t}(\mu_t) + \mathcal{A}^{\nu, \lambda, \sigma} V_{T-t}(\mu_t)] dt \right] \\
 &\geq \mathbb{E} [V_{T-\tau \wedge T}(\mu_{\tau \wedge T})] \\
 &\geq \mathbb{E} [\rho_T(\tau)]
 \end{aligned}$$

where the first equality uses Itô's formula for jump processes. It follows from $\partial V_{T-t}(\mu_t)/\partial \mu_t$ and $\partial V_{T-t}(\mu_t)/\partial \lambda_t^i$ are bounded which implies that the diffusion and jump terms in the drift of $V_{T-t}(\mu_t)$ are bounded. Since $\partial V/\partial T = U$ as noted in Step 2, the first inequality follows. The second inequality follows from the definition of V .

Step 3: $\{V_{T-t}(\mu_t^G)\}$ is a Martingale. It remains to verify that it satisfies the conditions of Lemma 3. We begin with Lemma 3.1. The inner and outer max on the right-hand side of (5) are satisfied. Exploitation then (5) is satisfied.

LEMMA 4 At $t \in [0, \infty)$ the following hold:

1) If $\mu \geq \mu^*$, then

$$U_t(\mu) = \frac{dV_t(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)}.$$

2) If $\mu \leq \mu^*$, then

$$U_t(\mu) = \frac{dV_t(\underline{\mu}, \mu)}{d_H(\underline{\mu}, \mu)}.$$

PROOF OF LEMMA 4:

Because $V_{T-t}(\mu_t^G) = \mathbb{E}[\rho_T(\tau_{\mu^G}) | \mu_t^G]$ and μ^G is Markov it follows that $\{V_{T-t}(\mu_t^G)\}$ is a martingale. For Bayly to give formula (17) the drift of $V_{T-t}(\mu_t^G)$ is zero if and only if conditions 1 and 2 of the lemma are satisfied.

⁸See Theorem 51 of Protter (2005).

Step 5. Unimprovability - By Lemma 5 shows that Greedy Exploitation cannot be improved on by any algorithm.

LEMMA 5 At $(\mu, \bar{\mu}) \in (\underline{\mu}, \bar{\mu}) \times [0, \infty)$ it holds that

$$(7) \quad U_t(\mu) = \max_{\nu} \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}.$$

PROOF OF LEMMA 5:

We will prove the lemma where $\mu \leq \bar{\mu}$ is analogous. By Lemma 4, it suffices to show that \max in (7). We split proof into three cases.

- Case $\mu \geq \bar{\mu}$. We will show that the global maximizer of $d_{V_t}(\nu, \mu)/d_H(\nu, \mu)$ in the region $\nu \geq \bar{\mu}$. To start, we observe that

$$\frac{d}{d\nu} \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)} = \frac{V'_t(\nu) - V'_t(\mu)}{d_H(\nu, \mu)} - \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)^2} [H'(\nu) - H'(\mu)].$$

This derivative is negative if and only if

$$(8) \quad \frac{V'_t(\nu) - V'_t(\mu)}{H'(\nu) - H'(\mu)} \geq \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)},$$

which is equivalent to

$$(9) \quad \frac{d_{V_t}(\bar{\mu}, \mu) - d_{V_t}(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \geq \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}.$$

Notice that (9) holds with equality at $\nu = \bar{\mu}$. We will show that in fact any local extremum of the function in (9) is a local maximum. This immediately implies the global maximum in this region.

At any local extremum (9) holds with equality. We local extrema are necessarily local maxima since the derivative of the left-hand side of (9) is negative. The right-hand side is always zero at a local extremum. Therefore the left-hand side of (9)

is decreasing

$$\begin{aligned} \frac{d}{d\nu} \frac{dV_t(\bar{\mu}, \mu) - dV_t(\bar{\mu}, \nu)}{dH(\bar{\mu}, \mu) - dH(\bar{\mu}, \nu)} &= \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\nu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &< \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\mu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &= 0 \end{aligned}$$

where we have used $dV_t(\bar{\mu}, \nu)/dH(\bar{\mu}, \nu) = U_t(\nu)$ from Lemma 4 and that this is increasing in Step 2.

- Case $\nu \notin (\mu^*, \mu)$. In this region, (following the same steps) it is easy to see that $dV_t(\nu, \mu)/dH(\nu, \mu)$ is nondecreasing if

$$(10) \quad \frac{dV_t(\bar{\mu}, \mu) - dV_t(\bar{\mu}, \nu)}{dH(\bar{\mu}, \mu) - dH(\bar{\mu}, \nu)} \leq \frac{dV_t(\nu, \mu)}{dH(\nu, \mu)}.$$

This is the same condition as (9) except the inequality

As before, to determine whether a local extremum exists, it suffices to check how the left-hand side of (10) changes as the left-hand side is increasing. This can be seen

$$\begin{aligned} \frac{d}{d\nu} \frac{dV_t(\bar{\mu}, \mu) - dV_t(\bar{\mu}, \nu)}{dH(\bar{\mu}, \mu) - dH(\bar{\mu}, \nu)} &= \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\nu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &> \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\mu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &= 0 \end{aligned}$$

where we have used the fact that the denominator is increasing. Thus, in this region, any local extremum must be attained at a point $\mu \in (\mu^*, \mu)$ can achieve the maximum in (7).

- Case $\nu \in [\mu, \mu^*]$. Following analogous steps to those used above, we find that $dV_t(\mu, \mu)/dH(\mu, \mu)$ is decreasing and only if

$$(11) \quad \frac{dV_t(\underline{\mu}, \mu) - dV_t(\underline{\mu}, \nu)}{dH(\underline{\mu}, \mu) - dH(\underline{\mu}, \nu)} < \frac{dV_t(\nu, \mu)}{dH(\nu, \mu)}.$$

We will prove that the left-hand side of (11) is bounded above by $dV_t(\nu, \mu)/dH(\nu, \mu)$. Thus, there can be no $\mu \in [\mu, \mu^*]$ that achieves a higher value of $dV_t(\mu, \mu)/dH(\mu, \mu)$, since if there were, it would not be decreasing in

To show this, we first observe that

$$(12) \quad d_{V_t}(\underline{\mu}, \mu) = d_{V_t}(\underline{\mu}, \bar{\mu}) + d_{V_t}(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (V'_t(\mu) - V'_t(\bar{\mu})),$$

and

$$(13) \quad d_H(\underline{\mu}, \mu) = d_H(\underline{\mu}, \bar{\mu}) + d_H(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (H'(\mu) - H'(\bar{\mu})).$$

Define $f(\mu)$ and $g(\mu)$ as

$$(14) \quad f(\mu) = d_{V_t}(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (V'_t(\mu) - V'_t(\bar{\mu}))$$

and

$$(15) \quad g(\mu) = d_H(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (H'(\mu) - H'(\bar{\mu})).$$

Since (8) holds, it follows that

$$(16) \quad \frac{f(\mu)}{g(\mu)} = \frac{d_{V_t}(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (V'_t(\mu) - V'_t(\bar{\mu}))}{d_H(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (H'(\mu) - H'(\bar{\mu}))} = \frac{d_{V_t}(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)}.$$

Also, $d_{V_t}(\bar{\mu}, \mu^*)/d_H(\bar{\mu}, \mu^*) = d_{V_t}(\underline{\mu}, \mu^*)/d_H(\underline{\mu}, \mu^*)$,

$$(17) \quad \frac{f(\mu^*)}{g(\mu^*)} = \frac{d_{V_t}(\underline{\mu}, \bar{\mu}) + f(\mu^*)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu^*)} \Rightarrow \frac{f(\mu^*)}{g(\mu^*)} = \frac{d_{V_t}(\underline{\mu}, \bar{\mu})}{d_H(\underline{\mu}, \bar{\mu})}.$$

Thus,

$$\begin{aligned} \frac{d_{V_t}(\underline{\mu}, \mu) - d_{V_t}(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \mu) - d_H(\underline{\mu}, \nu)} &= \frac{d_{V_t}(\underline{\mu}, \bar{\mu}) + f(\mu) - d_{V_t}(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu) - d_H(\underline{\mu}, \nu)} \\ &= \frac{U_t(\mu^*)d_H(\underline{\mu}, \bar{\mu}) + U_t(\mu)g(\mu) - U_t(\nu)d_H(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu) - d_H(\underline{\mu}, \nu)} \\ &\leq \frac{U_t(\mu^*)d_H(\underline{\mu}, \bar{\mu}) + U_t(\mu)g(\mu) - U_t(\mu^*)d_H(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu) - d_H(\underline{\mu}, \nu)} \\ &\leq U_t(\mu) \\ &= \frac{d_{V_t}(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)}. \end{aligned}$$

as desired. The first line uses (12), (13), (14), and (16), (17), and Lemma 4. The third line uses the fact that $U_t(\mu^*) \geq U_t(\mu)$ for $\mu \in [\underline{\mu}, \mu^*]$ as noted in Step 2. The last line uses Lemma

Step 6. Unimprovable by Definition 6 shows that Greedy Exploitation cannot be improved on by d

LEMMA 6 At $(\bar{\mu}, \mu) \in (\underline{\mu}, \bar{\mu}) \times [0, \infty)$, it holds that

$$U_t(\mu) \geq \frac{V_t''(\mu)}{H''(\mu)}.$$

PROOF:

Recall from Step 2 that $U_t(\mu) > 0$ whenever $(\mu, \bar{\mu}) \in (\underline{\mu}, \bar{\mu}) \times [0, \infty)$. It thus

$$U_t'(\mu) = \frac{d}{d\mu} \frac{dV_t(\bar{\mu}, \mu)}{dH(\bar{\mu}, \mu)} = \frac{-d_H(\bar{\mu}, \mu)V_t''(\mu)(\bar{\mu} - \mu) + dV_t(\bar{\mu}, \mu)H''(\mu)(\bar{\mu} - \mu)}{d_H(\bar{\mu}, \mu)^2} > 0$$

which implies that

$$U_t(\mu) = \frac{dV_t(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)} > \frac{V_t''(\mu)}{H''(\mu)}$$

as desired. An analogous argument We ask the question why when $\mu = \mu^*$ follows from continuity.

Step 7. Putting Lemma 5 together imply that (5) is satisfied for $\bar{\mu} \neq \mu^*$. Lemma 3 then implies that Greedy Exploitation any discount function ρ_t to $\rho_t = 0$ if $\bar{\mu} \neq \mu^*$ then implies optimality for all ρ . The proof of Theorem 1 is complete.

B. Time-Risk Averse

When the agent is time-risk averse, Pure Accumulation is optimal. It is illustrated graphically below in Figure 2.

As discussed in Section 1, the Pure Accumulation strategy is the suspense-maximizing strategy, a fact already established by Kamenica (2015). In this strategy, the agent's belief either jumps in the direction of experiences compensating drift. When her belief is with the same entropy as her current belief so that a drift.

DEFINITION 3: Pure Accumulation strategy is defined $\mu^H: [0, 1] \setminus \{\mu^*\} \rightarrow [0, 1]$ denote the function that maps a belief $\hat{\mu}$ to $\mu^H(\hat{\mu}) \neq \hat{\mu}$ such that $H(\mu^H(\hat{\mu})) = H(\hat{\mu})$. Under Pure Accumulation, the agent's beliefs evolve according to

$$d\mu_t^P = [\mu^H(\mu_t^P) - \mu_t^P] dJ_t(\lambda_t) - \lambda_t [\mu^H(\mu_t^P) - \mu_t^P] dt$$

where J_t is a Poisson point process with jumps at rate $\lambda_t = 1/S_H(\mu_t^H, \mu_t^A)$.

THEOREM 2 If the agent is time-risk averse, then Pure Accumulation is optimal.

PROOF:




Figure 2. Pure Accumulation

Note the dark green curve represents μ^B on the possible cable lengths of represent jumps. The light green curve represents μ^E on the other possible figure is computed $H(\hat{\mu}) = \tilde{\mu}^{2\gamma}$ the case

Under Pure Accumulation, the agent is guaranteed deterministic time. Because Pure Accumulation is eximmediately from Lemma 1. Because Pure Accumulation entails no time risk, we iing result.

COROLLARY 2 It holds that for each $\tau \in T$

V. Concluding Discussion

In this paper, we have studied the relationship between time risk and optimal information acquisition. We have shown that for a time-risk loving agent is Greedy Exploitation or riskiest distribution over threshold hitting times. On the other hand, an optimal strategy for a time-risk averse agent is Accumulation. This strategy produces a deterministic distribution and thus entails no time risk. Both of these strategies depend on the convexity or concavity of the utility function. Thus, they are immune to dynamic inconsistency. Thus, they may have time preferences that differ from the well-known exponential discounting. Our analysis offers some insight into how to model agents who acquire information and the kinds of signal structures that are optimal when modeling these agents.

In order to illustrate the connection between learning as sharply as possible we have made a number of spec

sumptions of binary states, fixed stopping threshold on learning speed are critical because they ensure have the same expected threshold-hitting times. ⁹ The reason why Greedy Exploitation and Pure Accumulation are optimal is because the maximal and minimal threshold-hitting times in order among all exhaustive strategies. This allows risk preferences that determines their optimality. They depend only on threshold hitting times and not on what we use to derive closed-form solutions but is not critical.

Indeed, we anticipate that many of the qualitative properties of Greedy Exploitation or Pure Accumulation will persist under more general settings. For example, with other costs of learning, multiple stopping thresholds (there are already examples in the literature as reviewed in Section I). It is certainly possible to generalize our results to accommodate these more general environments. Our characterization of optimal strategies is a characterization of our special setup. It is possible to solve for optimal strategies that are optimal for large classes of payoffs and, moreover, allows us to characterize optimal learning.

There are two promising avenues to explore in future work. First, we would like to explore how our results may extend to the case when the agent is risk loving or time-risk averse. For these more general settings, the qualitative features of optimal information acquisition and the qualitative features of optimal information acquisition to explore is to try to embed our model of information acquisition in settings where there are multiple agents in order to study flexible information acquisition in games.

REFERENCES

- Arrow, Kenneth J., David Blackwell, 1949, "Bayesian Epistemic Decision Making and Minimax Solutions of Sequential Decision Problems," *Journal of the Economic Theory* 1: 3-24.
- Caplin, Andrew, Mark Dean, 2017, "Rationality and Inattentive Behavior: Characterizing and Generalizing Sequential Decision Making," Bureau of Economic Research Working Paper 23652.

⁹Specifically, with a binary state and fixed thresholds, every learning strategy yields the same expected threshold-hitting times. That is, all learning strategies yield the same "over the long run" expected threshold-hitting times.

¹⁰Our solutions were based on a guess and verify approach that was possible due to the structure of our setup. In more general setups guessing the optimal strategy is more difficult.

- Chen, M Keith. 2013. "The effect of language on economic
 dence from savings rates, health ambiventions, and r
 Economic, R 103 (e2v) : 690-731.
- Chesson, Harrell W, and 2003. "VC commonalities in time
 ambiguity aversion Theedroyngnd, D 14 (15) : 657-71.
- Che, Yeon-Koo, and Konrad M. 2019. "Opportimal Dynamic All
 cation of AAternitdrom Econpmio 9 R 8 y i 2v 93-3029.
- Dasgupta, Partha, and 1980. "The effect of uncertainty, industri
 ture, and the s TheedBof R 8 D.urnal of 2Economics
- DeJarnette, Patrick, David Dillenberger, Daniel G
 Ortol 2020. "Time lotteries and econmstii caitpati
 88(2) : 619-656.
- Ely, Jeffrey, Alexander Franke 2015. "ESuisrpk an m e nai ncca .
 surpriosu.r.nal of Pol i ti 2c3a(11) : c 0n1o5my 260.
- Frankel, Alexander, and 2019. K 2manit c a y i n g i n f o r m a
 and unce r Atm e i r r i t c y a . n Econpmio 9 R 4 0) e 3 650-80.
- Fudenberg, Drew, Philipp Strack, 2018. "To s p a s e d, s t r z a l
 accuracy, and the opti m a t e r i i n i a m g E c o n o m i i c R e . v " i e w
 108(12) : 3651-84.
- Georgiadis-Ha 2023., "A P k e p a r i n g t o A c t . " w o r k i n g p a p
 Gossner, Olivier, Jakub Stein 2021 and A C d l e i n t i S t r e w a n
 p l e a E c o n o m e t r i c 4a) : 1717-1751.
- Hébert, Benjamin, and Mi 2017. "W a t d i f o r a d l i n a t t e n t i o
 sequential information sampling." National Burea
- Hébert, Benjamin, and Mi 2023. "W a d f o n d l i n a t t e n t
 when decision s o u r n e l t i o n e E c o n d m i 5 6 T h e o r y
- Hörner, Johannes, and An 2017. "S k e a y p a n g , E x p e r i m e
 tation, and I n f o A n a t h o n s D e s i E c n o . n o m i c s and Econome
 1: 63-98.
- Karatzas, Ioannis, and 1998. "B e r s h w n B r o . m o t i o n . "
 ian motion and stop 47a-s12i7c calculus
- Ke, T Tony, and J M i g u e l 2019. "I a s o p t i d a m a l l e a r n i n g b e t
 choi J c o u r n a l o f E c o n d m i o c B 18 3 e r 4 3 7.
- Lee, Tom, and Louis 1980. "M a r k e t s t r u c t u r e and i n n o v
 reformu T h e i Q u a r . t e r l y J o u r n 24 (o2f) E c 4 o 2 n 9 o - n i 3 c 6 s.

- Liang, Annie, Xiaosheng Mu, and David S. Foray. 2019. "Dynamics of aggregating dividers in a network." *Journal of Economic Theory* 191: 1910-07015.
- Loewenstein, George, and David Laibson. 1992. "Time discounting and time preference." *Journal of Economic Literature* 30(2): 573-597.
- Matijka, Filip, and Alisdair R. Meade. 2015. "Inattention to choices: A new foundation for the multinomial logit model." *Econometrica* 83(1): 272-298.
- Mayskaya, Tatyana. 2019. "Dynamic Choice of Information." *Working Paper*.
- Morris, Stephen, and Philip Tetlock. 2011. "The Problem and the Solution: Equivalence of Sequential Sampling and Static Inference." *Journal of Experimental Psychology: Applied* 17(1): 1-14.
- Moscarini, Giuseppe, and David Laibson. 2001. "The small level of exponential discounting." *Journal of Economic Theory* 91(1): 1629-1644.
- Müller, Albrecht. 1996. "Orderings of risks: A comparative statics approach." *Mathematical Economics* 24(1): 1-22.
- Nikandrova, Arina, and David Laibson. 2018. "Dynamic project selection." *Theoretical Economics* 13(1): 1-43.
- Onay, Selçuk, and Albrecht Müller. 2007. "On the temporal choice under risk: An experimental study." *Journal of Economic Theory* 121(1): 1-21.
- Protter, Philip. 2005. "Stochastic differential equations." *Journal of Mathematical Analysis and Applications* 304(1): 1-14.
- Ratcliff, Roger, and Jeffrey N. Rouder. 1998. "Modeling response time distributions." *Psychological Review* 105(1): 347-356.
- Reinganum, Jennifer. 1989. "The timing of innovation: Research, development, and diffusion." *Journal of Industrial Organization* 17(3): 347-367.
- Sims, Christopher. 2003. "Implications of rational inattention for monetary policy." *Journal of Monetary Economics* 51(4): 665-690.
- Steiner, Jakob, Colin Stewart, and Filip Matijka. 2017. "Inattention Dynamics: Inertia and Delay in Decision Making." *Journal of Economic Theory* 168: 85(2): 521-553.
- Wald, Abraham. 1947. "Foundations of a general theory of statistical estimation." *Econometrica* 15(1): 1-10.

Zhong, Weiye. "Optimal dynamic information acquisition in a continuous time framework." *Journal of Economic Theory* 90(4): 1537-1582.

Zhong, Weiye, and Alexander B. Wernz. "The cost of optimally acquired information." working paper.