

I n f o r m a t i o n A c q u i s i t i o n a n d T i m e - R i s k P r e f e r e n c e

D a n i e l a O d t c h e n⁰ W e i j i e Z h o n g[†]

S e p t e m b e r 7 , 2 0 2 4

A b s t r a c t

An agent acquires information dynamically until her belief state reaches an upper or lower threshold. She can choose to be subject to a constraint on the rate of entropy reduction by “time risk”—the dispersion of the distribution of beliefs. We construct a strategy G_{\max} and an E_k pointwise risk (minimizing Purgel Accumulation) for either strategy, belief state, and compensated Poisson process. In the former, belief state is closer in Bregman divergence. In the latter, belief state is closer with the same entropy as the current belief.

1 I n t r o d u c t i o n

In this paper, we study information acquisition by an agent about a binary state that may be either zero or one. The agent chooses to be subject to a constraint on the rate of entropy reduction by “time risk”—the dispersion of the distribution of beliefs. She earns a utility from learning but has great flexibility in how she can learn but has a constraint on her rate of learning. That is, she can choose any pointwise risk constraint on the rate of entropy reduction.

⁰odtchen@princeton.edu, Department of Economics, Princeton University
[†]wjzhong@stanford.edu, Graduate School of Business, Stanford University
We thank Kayla Jiang for excellent research assistance.

Our simple model captures three important features: flexible learning, limited resources, and thresholds. These appear in the contexts of research and development, marketing, user-experience testing, and others. For example, Facebook who must assess whether to introduce a new feature, an unknown state is whether adding the feature increases revenue (and thus profits). The data scientist can learn about the state by running tests. To provide incentives, her manager offers her a reward that is sufficiently precisely about the state (i.e. her manager offers a minimal level of¹) such that she can sustain a profitable design many aspects of the design—e.g., she can select the subpopulation of users to test the feature and she can adjust the length of time a user has to make other choices. However, there are limits to what she can learn—e.g., her manager does not allow her to run multiple tests simultaneously as doing so could be disastrous for the company. She can only implement the feature for a given user.

Our main contribution is to show how, in such settings, a learning strategy depends on risk preferences, costs, and discounting. We say that the agent is time-risk neutral if her utility over threshold-hitting times beyond which the feature is discounted is linear. We show that the optimal strategy is time-risk neutral whenever she is time-risk averse. Critically, the optimal strategy depends on the shape of the agent's utility function.

In reality, there are many reasons why individuals differ from the predominantly studied case of expected utility. These include external factors such as explicit discounting, flow costs associated with foregone opportunities, and internal factors such as present bias resulting from a simple framework that allows for these factors.

¹In our binary state model, there is a one-to-one mapping between the power of a test (that is, the likelihood of Type I errors) and the reward.

²In this paper, we model the agent as an expected-utility maximizer. That all of our results will go through as long as the agent's utility function is monotonic in the mean-preserving spread order.

derive strategies that are uniformly optimal up to utility function over threshold-hitting times.

We now briefly describe the two learning strategies. When the agent is **time-risk averse**, **Greedy Exploitation** is optimal. In this strategy, the agent myopically maximizes the jump to a threshold. She acquires a rare but decisive belief to jump to the threshold that induces her belief to jump to the threshold that is closest to the target. By targeting the closer threshold, she can jump at a lower rate, relaxing her constraint on the rate of entropy reduction. As her belief experiences compensating drift in the direction of the target, eventually, her belief reaches a point that is equidistant from the two thresholds. At this point, she acquires a belief that induces her to jump to either threshold but at rates set so that her belief is stationary in the absence of new information.

Intuitively, Greedy Exploitation is optimal because it minimizes the contribution of threshold-hitting times to the overall risk. Because the probability of an early hitting time is higher, the agent jumps earlier. As beliefs drift towards the farther threshold, the jump to the closer threshold is delayed, and the amount of time remaining until the jump is reduced. In this way, the agent makes no "progress" in the absence of a jump to the closer threshold, and late threshold hitting times as well. We, in fact, show that this exhausts the agent's resources (in that the constraint on entropy reduction is binding at all points in time), Greedy Exploitation achieves threshold hitting times that is maximal in the mean-preserving sense, and thus minimizes time risk.

When the agent is **time-risk neutral**, she instead uses **Pure Accumulation**. In this case, an optimal strategy is to reach a threshold as late as possible. Her beliefs follow a compensating process that jumps in the direction of the threshold as needed to maintain an interior belief that has the same entropy as her belief after a jump. Pure accumulation is a continuous-time analog of the strategy in Ely et al. (2015). The strategy is, in effect, the same as Greedy Exploitation. Because jumps are always to beliefs with the same entropy, the agent's belief is stationary in the absence of new information.

event of a jump, there is no progress: a jump does not stop time until a threshold is hit. Instead, all progress is made before the jump. This is why the threshold hitting time is deterministic. It is not a time risk. It, therefore, produces a distribution of hitting times with a mean-preserving spread over all strategies (in that the constraint on the rate of entropy reduction is a time risk).

Our analysis of optimal learning through the implications for both information acquisition in the model predicts that an agent who is time-risk-loving, whereas an agent who is time-risk-averse is not optimal, provided the agent has access to these learning strategies. The agent's space of available learning strategies is generally suboptimal. Thus, when writing models with parameterized signal structures, economists can assume signal structures are without loss of optimality for the agents they seek to model.

2 Related Literature

Our paper contributes to a large literature on information acquisition (1947) and Arrow et al. (1949) we study a sequential learning problem where the agent to flexibly design the signal process as in Crawford (2017), Hébert & Woodford (2023), Steiner et al. (2023). Whereas most of these papers restrict attention to exponential discounting or a linear delay cost, we allow for a general discounting function. For example, Zhong (2022) assumes exponential discounting and time-risk-loving preferences whereas Hébert & Woodford (2023) assumes time-risk-averse preferences and a linear delay cost which implies time-risk-loving preferences. We suggest that the assumed time-risk preferences do not affect the optimal strategies identified in these papers.

³Hébert & Woodford (2023) allow for both discounting and a linear delay cost. We consider the time-risk neutral limit for the majority of the papers. We study how different costs or constraints on information acquisition affect the optimal strategies which is orthogonal to the objective of our paper. Zhong (2022) studies the time-risk loving case.

Pure Accumulation is a continuous-time variant introduced by Ely et al. (2015) in a discrete-time time horizon. In Ely et al. (2015), the suspense- it maximizes an expected utility that is increasing. Georgiadis-Harris (2023) also finds that a strategy optimal in a setting where the agent does not depend on the learning time lags is behind the optimality of in both of these papers are distinct from that of time-risk aversion.

In our analysis, the key summary statistic that strategy is the distribution of the time that the agent. This statistic is an object studied in a literature on time-risk preferences. Chesson & Viscusi the expected discounted utility framework is improved over time (Rust). DesJardette et al. (2020) show that of models RSTL cannot be violated if there is stopping experimental evidence suggests that a better alternative (RATL) (Chesson & Viscusi (2003); Onay & Öncüler modulates both RSTL and RATL and shows that optimal differ dramatically under different time-risk preferences.

The optimal learning strategies that we identify learning strategies that have been assumed in reduced form literature. For example, Che & Miendorff (2019), Mays & Pans (2018) adopt a framework that restricts attention in order to study optimal stopping with endogenous also often assumed in the literature on strategic Hörner & Skrzypacz (2017)). We show that Poisson foundation under time ⁵ if the knowledge of the environment is also related to classic models on the timing of information function that does not have a threshold structure but show nevertheless similar to Greedy Exploitation.

⁴In our paper, the stopping is determined by the chosen learning strategy.

⁵To be clear, we do not show that the learning strategies are optimal strategies in our setting involve Poisson learning.

Stiglitz (1980) and Lee & Wilde (1980) (see a survey in Smith (2001)). They involve a deterministic time of innovation. The reduced-form learning process and are non-Bayesian learning strategies in these papers can emerge in an information acquisition framework when agents have time costs.

Our model also allows for Gaussian learning strategies, which are often assumed in reduced-form learning models (e.g., Smith (2001); Ke & Villas-Boas (2019); Liang et al. (2021)). Although, if π is a binary choice problem, applying the Rousseeuw (1998) and Fudenberg et al. (2018) model, Bayesian learning cannot be justified by optimality except if agents have time-risk neutral preferences provided information is costly.

The optimality of a greedy strategy is also the main result of our paper. However, the mechanisms in our papers are very different from the previous ones. Crucially, it depends on the linear-Gaussian setup with multiple information sources and holds for any time preferences and endogenous choice of information sources, but not for exogenous ones. Also related is Gossner et al. (2021) which considers a stopping time (when the belief hits a threshold) in the sense of first-order stochastic dominance.

We model limits on the agent's learning resource by the entropy reduction. That is, the rate of information acquisition is bounded. The rational inattention (RI) typically models information costs or constraint on the number of signals (Matzka & McKay (2014); Steiner et al. (2017); Caramazza et al. (2017); Zhong & Bloedel (2021); Morris & Strzalecki (2021)). Information constraint ensures that the expected information gain is bounded for all exhaustive strategies, which is a key property for information acquisition. By Theorem 3 in Zhong (2022), the boundedness is both necessary and sufficient for the expected learning gain to be bounded for exhaustive strategies.

This section presents a simple model of an agent with an unknown state. The unknown state $\mu \in (0, 1)$ is set at $t = 0$, the agent believes it has the prior $\mu \in (0, 1)$ and receives a unit payoff μ at time $t = 1$. The agent's belief μ_t evolves as a martingale and reaches either $\bar{\mu} \in (0, 1)$ or $\underline{\mu} \in (0, \mu)$ at time $t = 1$. However, she is impatient and her utility is $\rho: \mathbb{R}_+ \rightarrow \mathbb{R}$ of the threshold-hitting time τ where $\rho(t) = e^{-rt}$. On the other hand, she is also aware of the cost of delay that increases exponentially with time:

The agent has great flexibility in using her resources and cannot learn about the state of the world. She has a sequence of beliefs $\{\mu_t; t \geq 0\}$ that take values in $[0, 1]$ and satisfy a stochastic differential equation (allowing for jumps) of the form

$$d\mu_t = \sum_{i=1}^N (\nu^i(t, \mu_t) - \mu_t) [dJ_t^i(\lambda^i(t, \mu_t)) - \lambda^i(t, \mu_t) dt] + \sum_{j=1}^M \sigma^j(t, \mu_t) dZ_t^j \quad (1)$$

with $\mu_0 = \mu$ for some positive μ and ν^i, σ^j are N, M functions of μ_t and Z_t^j is a standard Brownian motion. The process μ_t is a point process that takes values in $[0, 1]$ and the number of distinct points that the process visits is the number of distinct Brownian Motions.

We assume that the agent can direct her search to any

$$\mathbb{E} [H(\mu_s) - H(\mu_t) | \mathcal{F}_t] \leq I(s - t) \quad (2)$$

for t, s such that $t < s$. The \mathcal{F}_t is the natural filtration of μ_t . If H is a strictly concave function defined on $[0, 1]$, then the agent's belief μ_t is a martingale. Equation (2) is a constraint on the reduction. A special case is when H is the Shannon entropy so that (2)

⁶Our restriction to jump-diffusion belief processes is within the class of càdlàg processes such that (2) is well defined. This is Harris (2023).

a constraint on the well-known mutual information we normalize $H(\underline{\mu}) = H(\bar{\mu}) = 0$ and $d = 1$. This can be done by redefining

$$\frac{1}{I} \left[H(\mu) - \frac{H(\bar{\mu}) - H(\underline{\mu})}{\bar{\mu} - \underline{\mu}} (\mu - \underline{\mu}) \right].$$

The same belief processes satisfy (2) before and after the normalization and are martingales (and thus, the drift of the second term is zero). The normalization is convenient because, provided optional stopping is ethically acceptable, the expected time remaining until the belief process hits the boundary is simply the current entropy:

$$-H(\mu_t) = \mathbb{E}[H(\mu_\tau) - H(\mu_t) | \mathcal{F}_t] = \mathbb{E}[\tau - t | \mathcal{F}_t].$$

To state the agent's belief threshold, let τ_μ be the first time the belief process hits the boundary:

$$\tau_\mu := \inf\{t | \mu_t \in [0, \underline{\mu}] \cup [\bar{\mu}, 1]\}.$$

Because μ may be with positive probability for some belief, the agent can stop (or not) at any time. We assume that the agent never stops these processes.

She solves

$$\max_{\mu \in \mathcal{M}} \mathbb{E}[\rho(\tau_\mu)] \quad (3)$$

such that (2) holds.

Our simple model makes several assumptions in order to bridge between optimal learning and time-risk preferences in the paper. For example, we assume that the agent expects to live long as (2) is satisfied and that she earns a common return on capital that she ultimately hits. Though these assumptions are simplifying, our model aligns well with a number of economic assumptions. Introduction. A key property of our model is that it accommodates general time-risk preferences. We identify optimal learning strategies that apply

⁷See Theorem 3.22 of Karatzas & Shreve (1998).

⁸That is, we extend the domain and range of

and to highlight that it is the agent's decision that determines the qualitative properties of these strategies. preferences beyond exponential discounting include $\rho(t) = (1 + \alpha t)^{-\gamma/\alpha}$ (Loewenstein & Prelec, 1992); (ii) general utility $u(x) = \phi(e^{-rx})$ where ϕ is increasing and convex (DeJonghe & linear delay cost up to T and 0 thereafter (risk averse preferences are less commonly assumed in the literature experimental evidence that agents may often be time 2003; Onay & Öncüler, 2007). As discussed at the aversion can arise when agents have low costs of d

4 Optimal Learning and Time-Risk

In this section, we present our main results: a strategy that the agent is time-risk loving and a strategy that is time-risk averse. These results illustrate the connection between time preferences.

4.1 Time-Risk Loving

We first consider the case when the agent is time-risk loving. Her optimal learning strategy is given below in Definition 1 informally here.

An optimal strategy for a time-risk loving agent is illustrated in Figure 1.

Let

$$d_H(\tilde{\mu}, \hat{\mu}) = H(\tilde{\mu}) - H(\hat{\mu}) - H'(\hat{\mu})(\tilde{\mu} - \hat{\mu})$$

denote the Bregman divergence between $\tilde{\mu}$ and $\hat{\mu}$ relative to H . In Figure 1, μ^* represents the unique belief that is equal to the divergence to the threshold κ . Initially, the agent either jumps to the threshold that is closest to her experience compensating drift towards the other

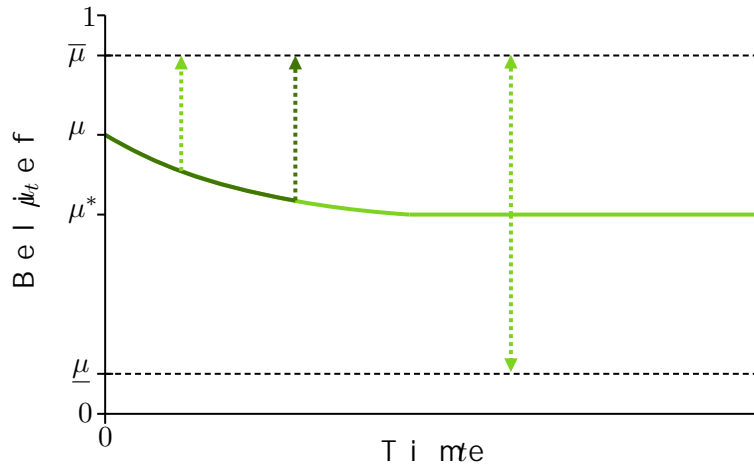


Figure 1: Greedy Exploitation

Notes: In dark green, we plot one possible cumulated realization of the Greedy Exploitation strategy. In light green, we plot one possible realization of the Greedy Exploitation strategy. The dashed lines with arrows represent jumps in the belief process. The horizontal axis is time, and the vertical axis is belief. The initial belief is μ . The thresholds are $\bar{\mu}$ and $\underline{\mu}$. The belief process evolves according to the SDE (1) when $\mu_t < \bar{\mu}$ and according to the SDE (2) when $\mu_t > \underline{\mu}$.

threshold, the agent greedily maximizes the “characteristic” in the “next instant.” This is because the agent’s belief evolves according to the SDE (1) when she targets the closer threshold without violating the threshold. After some time, in the absence of a jump, the belief process reaches a point, her beliefs may jump to either threshold. If the thresholds are such that there is no net compensation for a jump, her beliefs remain stationary.

Definition 6.1 Greedy Exploitation strategy. Let $\mu^* \in (0, 1)$ be the unique belief μ^* such that $d_H(\mu, \mu^*) = d_H(\mu, \bar{\mu})$, that

- When $\mu_t^G > \mu^*$, her beliefs evolve according to

$$d\mu_t^G = (\bar{\mu} - \mu_t^G) [dJ_t^1(\lambda_t) - \lambda_t dt]$$

where $\lambda_t \in I/d_H(\bar{\mu}, \mu_t^G)$.

- When $\mu_t^G = \mu^*$, her beliefs evolve according to

$$d\mu_t^G = (\bar{\mu} - \mu_t^G) dJ_t^2 \left(\frac{\mu_t^G - \underline{\mu}}{\bar{\mu} - \underline{\mu}} \lambda^* \right) + (\underline{\mu} - \mu_t^G) dJ_t^3 \left(\frac{\bar{\mu} - \mu_t^G}{\bar{\mu} - \underline{\mu}} \lambda^* \right)$$

where $\lambda^* = 1/d_H(\bar{\mu}, \mu^*)$.

- When $\mu_t^G < \mu^*$, her beliefs evolve according to

$$d\mu_t^G = (\underline{\mu} - \mu_t^G) [dJ_t^1(\lambda_t) - \lambda_t dt]$$

where $\lambda_t = 1/d_H(\underline{\mu}, \mu_t^G)$.

Above J_t^1 and J_t^2 , are independent Poisson point processes indicated in parentheses.

Theorem 1 If the agent is time-risk loving, then Greedy

Before we sketch the proof of Theorem 1, we note that the threshold is uniformly optimal for a distribution of the riskiest threshold hitting times among all strategies that are optimal at all points in time. To make this precise, we first

Definition 1 Let \mathcal{M} be a set of distributions μ such that (2) binds at all

The Greedy Exploitation strategy produces a threshold that is minimal in the mean-preserving spread order among all thresholds in the set.

Corollary 1 Let $\mu \in \mathcal{M}$ and $\tau \in \mathcal{T}$

This result hinges on our assumption that the coin is fair in (2). Because of this assumption, all thresholds are equally likely to hit the threshold hitting time, which is equal to the initial

Proof of Theorem 1 The proof proceeds in seven steps.

Step 1. Set of Basis Functions. We show that any nonnegative function ρ can be written as a conical combination of functions

$$\rho_T(t) = \max\{T - t, 0\}$$

where $T \geq 0$. Thus, if Greedy Exploitation is optimal, it must be optimal for any nonnegative dividend process. For any process, it may take on negative values.

Lemma 1 If Greedy Exploitation is optimal for $T \geq 0$, it is optimal for $T \geq 0$.

See Theorem 3.6 in Müller (1996).

Step 2. Candidate Value Function. Let $V(\mu, T)$ denote the value function for Greedy Exploitation for $T \geq 0$. Let

$$V(\mu, T) = \begin{cases} \int_0^T (T-t) \lambda_t^G e^{-\int_0^t \lambda_z^G dz} dt, & \mu \in (\underline{\mu}, \bar{\mu}) \\ T, & \mu \in \{\underline{\mu}, \bar{\mu}\}. \end{cases} \quad (4)$$

In what follows, it is $\partial V(\mu, T) / \partial \mu = U(\mu, T)$ whenever $\mu \in (\underline{\mu}, \bar{\mu})$ where

$$U(\mu, T) = \int_0^T \lambda_t^G e^{-\int_0^t \lambda_z^G dz} dt$$

is the probability that a jump occurs before time T . To show $\partial U(\mu, T) / \partial \mu > 0$ if $\mu \in [\mu^*, \bar{\mu})$ and $\partial U(\mu, T) / \partial \mu < 0$ if $\mu \in (\underline{\mu}, \mu^*]$.

To ease the exposition, we adopt $V_T(\mu) = V(\mu, T)$ and $U_T(\mu) = U(\mu, T)$. Also, given a function f , let

$$d_f(\nu, \mu) = f(\nu) - f(\mu) - f'(\mu)(\nu - \mu)$$

where $f'(\mu)$ is well-defined. Note that $d_f(\nu, \mu) \geq 0$ if and only if f is convex.

Step 3. Verify that Greedy Exploitation is optimal. Use the following Lemma 2 which states that the Hamilton-Jacobi-Bellman (HJB) equation (5).

⁹To apply Theorem 3.6 in Müller (1996) easily, it is necessary to show that the set of admissible strategies is exhaustive and thus has the expected threshold hitting time.

$T > 0$ $V_T(\mu) \geq 2[m_i(\mu) - \rho_T]$ a t i s f i e s

$$U_t(\mu) = \max_{\nu} \left\{ \frac{dV_t(\nu, \mu)}{d_H(\nu, \mu)}, \frac{V_t''(\mu)}{H''(\mu)} \right\} \quad (5)$$

a t $t \in (\underline{\mu}, \bar{\mu}) \times [0, T]$ i t h o l d s (3) w i t h $\rho_T = \rho^T$

We f i r s t a s s e r t t h a t c o n d i t i o n (5) i s e q u i v a l e n t t o

$$U_t(\mu) = \max_{\{\nu^i\}, \{\lambda^i\}, \sigma} \mathcal{A}^{\nu, \lambda, \sigma} V_t(\mu) \quad (6)$$

$$\text{s.t. } \mathcal{A}^{\nu, \lambda, \sigma} H(\mu_t) \leq 1$$

w h e r e $\mathcal{A}^{\nu, \lambda, \sigma}$ i s t h e o p e r a t o r d e f i n e d b y f u n c t i o n s

$$\mathcal{A}^{\nu, \lambda, \sigma} f(\mu) = \sum_i \lambda^i d_f(\nu^i, \mu) + \frac{1}{2} \sum_j (\sigma^j)^2 f''(\mu).$$

T h a t $\mathcal{A}^{\nu, \lambda, \sigma}$ i s t h e i n f i n i t e s i m a l g e n e r a t o r f o r t h e c o m p o s e d p r o c e s s (1) . B e c a u s e $\mathcal{A}^{\nu, \lambda, \sigma}$ i s a d d i t i v e l y s e p a r a b l e , i t s u f f i c e s t o p o i n t o r v o l a t i l i t y t o a c h i e v e t h e m a x i m u m i n (6) . T h e r e f o r e , t o m a x i m i z e t h e " b a n g - f o r - t h e - b u d g e t " o p t i m a l i t y i s t h e s a m e a s t o m a x i m i z e t h e e x p e c t e d u t i l i t y o f H . T h e r e f o r e , (5) a n d (6) m u s t b e e q u i v a l e n t .

N e x t , s u p p o s e t h a t (5) i s s a t i s f i e d . C o n s i d e r a n a r b i t r a r y p a t h μ_t w i t h i n d u c e d f r o m t h e w e b h a v e h i t t i n g t i m e

$$\begin{aligned} V_T(\mu) &= \mathbb{E} \left[V_{T-\tau \wedge T}(\mu_{\tau \wedge T}) - \int_0^{\tau \wedge T} [-U_{T-t}(\mu_t) + \mathcal{A}^{\nu, \lambda, \sigma} V_{T-t}(\mu_t)] dt \right. \\ &\quad + \sum_j \int_0^{\tau \wedge T} \frac{\partial V_{T-t}(\mu_t)}{\partial \mu} \sigma_t^j dZ_t^j \\ &\quad \left. + \sum_i \int_0^{\tau \wedge T} [V_{T-t}(\nu_t^i) - V_{T-t}(\mu_t)] (dJ_t^i(\lambda_t^i) - \lambda_t^i dt) \right] \\ &= \mathbb{E} \left[V_{T-\tau \wedge T}(\mu_{\tau \wedge T}) - \int_0^{\tau \wedge T} [-U_{T-t}(\mu_t) + \mathcal{A}^{\nu, \lambda, \sigma} V_{T-t}(\mu_t)] dt \right] \\ &\geq \mathbb{E} [V_{T-\tau \wedge T}(\mu_{\tau \wedge T})] \\ &\geq \mathbb{E} [\rho_T(\tau)] \end{aligned}$$

where the first equality uses Itô's formula for j , $\partial V/\partial T = U$ as noted in Step 2, the second equality for $\partial V_{T-t}(\mu_t)/\partial \mu$ and V_{T-t} are bounded which implies that the differentials are true martingales, the inequality follows from (6), follows from the definition of \square

Step 4. $\{V_{T-t}(\mu_t^G)\}_{t \in \mathcal{M}}$ — The remaining steps satisfy the conditions of Lemma 2. We begin with Lemma 3 and outer max on the right-hand side of (5) is achieved. (5) is satisfied.

Lemma 3. Let $\mu \in [0, \infty)$;

$$\mu \geq \mu^* \quad \mathcal{M}$$

$$U_t(\mu) = \frac{d_{V_t}(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)}.$$

$$\mu \leq \mu^* \quad \mathcal{M}$$

$$U_t(\mu) = \frac{d_{V_t}(\mu, \mu)}{d_H(\mu, \mu)}.$$

Because $V_{T-t}(\mu_t^G) = \mathbb{E}[\rho_T(\tau_{\mu^G}) | \mu_t^G]$ and μ^G is Markov it follows that $\{V_{T-t}(\mu_t^G)\}$ is a martingale. For Bayly's formula, the $\{V_{T-t}(\mu_t^G)\}$ is zero if and only if conditions 1 and 2 of the

Step 5. Unimprovable — The proof follows from Lemma 4 showing Greedy Exploitation can not be improved on by any strategy.

Lemma 4. Let $(\mu, \bar{\mu}) \in (\underline{\mu}, \bar{\mu}) \times [0, \infty)$;

$$U_t(\mu) = \max_{\nu} \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (7)$$

Proof of Lemma 4. We prove the case $\mu > \bar{\mu}$. The other case is analogous. By Lemma 3, $\bar{\mu}$ satisfies the maximum. We split the proof into three cases.

¹ See Theorem 51 of Protter (2005).

- Case $\nu \geq \mu$. We will show that the global derivative of $d_{V_t}(\nu, \mu)/d_H(\nu, \mu)$ in the region $\nu \geq \mu$ starts, we observe that

$$\frac{d}{d\nu} \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)} = \frac{V'_t(\nu) - V'_t(\mu)}{d_H(\nu, \mu)} - \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)^2} [H'(\nu) - H'(\mu)].$$

This derivative is negative if and only if

$$\frac{V'_t(\nu) - V'_t(\mu)}{H'(\nu) - H'(\mu)} \geq \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (8)$$

which is equivalent to

$$\frac{d_{V_t}(\bar{\mu}, \mu) - d_{V_t}(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \geq \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (9)$$

Notice that (9) holds if $d_{V_t}(\bar{\mu}, \mu) \geq d_{V_t}(\bar{\mu}, \nu)$ and $d_H(\bar{\mu}, \mu) \geq d_H(\bar{\mu}, \nu)$. We will show that any local extremum of $d_{V_t}(\nu, \mu)/d_H(\nu, \mu)$ in the region $\nu \geq \mu$ is a local maximum. This immediately implies that $\bar{\mu}$ must be a global maximum.

At any local extremum (9) holds with equality. If $\bar{\mu}$ is a local extremum, then the left-hand side of (9) is negative. This is because the left-hand side is always zero at a local extremum since $d_{V_t}(\bar{\mu}, \mu) = d_{V_t}(\bar{\mu}, \nu)$ and $d_H(\bar{\mu}, \mu) = d_H(\bar{\mu}, \nu)$. The left-hand side of (9) is zero because

$$\begin{aligned} \frac{d}{d\nu} \frac{d_{V_t}(\bar{\mu}, \mu) - d_{V_t}(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} &= \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\nu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &< \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\mu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &= 0 \end{aligned}$$

where we have used $d_{V_t}(\bar{\mu}, \mu) = U_t(\mu)d_H(\bar{\mu}, \mu)$ and $d_{V_t}(\bar{\mu}, \nu) = U_t(\nu)d_H(\bar{\mu}, \nu)$ from Lemma 3 and that $U_t(\nu)$ is increasing in Step 2.

- Case $\nu \leq \mu^*$. In this region, (following the same

easy to show that $d_{V_t}(\bar{\mu}, \mu)/d_H(\bar{\mu}, \mu)$ is nondecreasing if

$$\frac{d_{V_t}(\bar{\mu}, \mu) - d_{V_t}(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \leq \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (10)$$

This is the same condition as (9) except the inequality is reversed. As before, to determine whether a local extremum satisfies the condition, we check how the left-hand side is increasing. This can be seen

$$\begin{aligned} \frac{d}{d\nu} \frac{d_{V_t}(\bar{\mu}, \mu) - d_{V_t}(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} &= \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\nu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &> \frac{d}{d\nu} \frac{U_t(\mu)d_H(\bar{\mu}, \mu) - U_t(\mu)d_H(\bar{\mu}, \nu)}{d_H(\bar{\mu}, \mu) - d_H(\bar{\mu}, \nu)} \\ &= 0 \end{aligned}$$

where we have used the fact that the denominator in this region, any local extremum must be a local maximum. $\nu \in (\mu^*, \mu)$ can achieve the maximum in (7).

- Case $\nu \in [\mu, \mu^*]$. Following analogous steps to those in Case 1, we find that $d_{V_t}(\mu, \mu)/d_H(\mu, \mu)$ is decreasing only if

$$\frac{d_{V_t}(\underline{\mu}, \mu) - d_{V_t}(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \mu) - d_H(\underline{\mu}, \nu)} < \frac{d_{V_t}(\nu, \mu)}{d_H(\nu, \mu)}. \quad (11)$$

We will prove that the left-hand side of (11) is bounded above by $d_{V_t}(\bar{\mu}, \mu)/d_H(\bar{\mu}, \mu)$. Thus, there can not be a point μ that achieves a higher value of $d_{V_t}(\bar{\mu}, \mu)/d_H(\bar{\mu}, \mu)$, since if there was, it would be decreasing in μ .

To show this, we first observe that

$$d_{V_t}(\underline{\mu}, \mu) = d_{V_t}(\underline{\mu}, \bar{\mu}) + d_{V_t}(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu})(V'_t(\mu) - V'_t(\bar{\mu})), \quad (12)$$

and

$$d_H(\underline{\mu}, \mu) = d_H(\underline{\mu}, \bar{\mu}) + d_H(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu})(H'(\mu) - H'(\bar{\mu})). \quad (13)$$

Define $f(\mu)$ and $g(\mu)$ as

$$f(\mu) = d_{V_t}(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (V'_t(\mu) - V'_t(\bar{\mu})) \quad (14)$$

and

$$g(\mu) = d_H(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (H'(\mu) - H'(\bar{\mu})). \quad (15)$$

Since (8) holds, it follows that

$$\frac{f(\mu)}{g(\mu)} = \frac{d_{V_t}(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (V'_t(\mu) - V'_t(\bar{\mu}))}{d_H(\bar{\mu}, \mu) - (\underline{\mu} - \bar{\mu}) (H'(\mu) - H'(\bar{\mu}))} = \frac{d_{V_t}(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)}. \quad (16)$$

Also, $d_{V_t}(\bar{\mu}, \mu^*)/d_H(\bar{\mu}, \mu^*) = d_{V_t}(\underline{\mu}, \mu^*)/d_H(\underline{\mu}, \mu^*)$,

$$\frac{f(\mu^*)}{g(\mu^*)} = \frac{d_{V_t}(\underline{\mu}, \bar{\mu}) + f(\mu^*)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu^*)} \Rightarrow \frac{f(\mu^*)}{g(\mu^*)} = \frac{d_{V_t}(\underline{\mu}, \bar{\mu})}{d_H(\underline{\mu}, \bar{\mu})}. \quad (17)$$

Thus,

$$\begin{aligned} \frac{d_{V_t}(\underline{\mu}, \mu) - d_{V_t}(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \mu) - d_H(\underline{\mu}, \nu)} &= \frac{d_{V_t}(\underline{\mu}, \bar{\mu}) + f(\mu) - d_{V_t}(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu) - d_H(\underline{\mu}, \nu)} \\ &= \frac{U_t(\mu^*)d_H(\underline{\mu}, \bar{\mu}) + U_t(\mu)g(\mu) - U_t(\nu)d_H(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu) - d_H(\underline{\mu}, \nu)} \\ &\leq \frac{U_t(\mu^*)d_H(\underline{\mu}, \bar{\mu}) + U_t(\mu)g(\mu) - U_t(\mu^*)d_H(\underline{\mu}, \nu)}{d_H(\underline{\mu}, \bar{\mu}) + g(\mu) - d_H(\underline{\mu}, \nu)} \\ &\leq U_t(\mu) \\ &= \frac{d_{V_t}(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)}. \end{aligned}$$

as desired. The first line uses (12), (13), (14), (16), (17), and Lemma 3. The third line uses the fact that $U_t(\nu) \leq U_t(\mu)$ for $\nu \in [\underline{\mu}, \mu^*]$ as noted in Step 2. The last line uses Lemma 3.

□

Step 6. Unimprovable by the default linear program. Lemma 5 shows that Greedy Exploitation cannot be improved on by different algorithms.

$$(\mu, t) \in (\underline{\mu}, \bar{\mu}) \times [0, \infty) \text{ i. ii.}$$

$$U_t(\mu) \geq \frac{V_t''(\mu)}{H''(\mu)}.$$

Recall from Step 2 that $V_t'(\mu) > 0$ where $\mu \in (\underline{\mu}, \bar{\mu})$. Thus

$$U_t'(\mu) = \frac{d}{d\mu} \frac{dV_t(\bar{\mu}, \mu)}{dH(\bar{\mu}, \mu)} = \frac{-d_H(\bar{\mu}, \mu)V_t''(\mu)(\bar{\mu} - \mu) + d_{V_t}(\bar{\mu}, \mu)H''(\mu)(\bar{\mu} - \mu)}{d_H(\bar{\mu}, \mu)^2} > 0$$

which implies that

$$U_t(\mu) = \frac{dV_t(\bar{\mu}, \mu)}{d_H(\bar{\mu}, \mu)} > \frac{V_t''(\mu)}{H''(\mu)}$$

as desired. An analogous argument works when $\mu = \mu^*$ follows from continuity. \square

Step 7. Putting Lemma 4 together imply that (5) is true. Lemma 2 then implies that Greedy Exploitation discount function of Lemma 1 then implies optimal. The proof of Theorem 1 is complete. \square

4.2 Time - Risk Averse

When the agent is time-risk averse, the Accumulation illustrated graphically below in Figure 2.

As discussed in Section 2, the Pure Accumulation is the suspense-maximal strategy (2015). Under this belief either jumps in the direction of the farthest satiating drift. When her belief jumps, it jumps to her current belief so that all progress is made then.

Definition 3. Pure Accumulation strategy is defined on $[0, 1] \setminus \{\mu^*\}$ follows. $\{\mu^*\} \rightarrow [0, 1]$ denote the function that maps μ to $\hat{\mu}$ if such $H(\mu^H(\hat{\mu})) = H(\hat{\mu})$. Under Pure accumulation, the age




Figure 2: Pure Accumulation

Note the dark green curve represents μ^P on the possible belief segments represent jumps. The light green curve represents μ^H . The figure is composed of two segments for the case $H(\mu) = \mu$.

according to

$$d\mu_t^P = [\mu^H(\mu_t^P) - \mu_t^P] dJ_t(\lambda_t) - \lambda_t [\mu^H(\mu_t^P) - \mu_t^P] dt$$

where J_t is a Poisson point process with $\lambda_t = 1/\delta_H(\mu^H(\mu_t^P), \mu_t^P)$ ticks at rate

Theorem 2. If the agent is time-risk averse, then Pure

Proof. Under Pure Accumulation, the agent is guaranteed to determine $\mu^H(\mu)$ in finite time. \square

Because Pure Accumulation entails no time risk, the following result.

Corollary 2.01. It is that $\tau \in \mathcal{T}$

$\mathcal{M} \mathcal{M} \mathcal{M}$

In this paper, we have studied the relationship between optimal information acquisition. We have shown that

loving agent is Greedy Exploitation. This strategy over threshold hitting times among all exhaustive optimal strategy for a time-risk averse agent is produces a deterministic threshold hitting time. of these strategies are uniformly optimal up to the utility function, provided the agent is impatient. This inconsistency. In practice, agents may have time well-studied case of exponential discounting. How these agents may seek to acquire information economists may consider using when modeling these

In order to illustrate the connection between learning as sharply as possible we have made a number of specifications of binary states, fixed stopping thresholds, speed are critical because they ensure that all expected threshold hitting times are negative. ¹¹Why Greedy Exploitation and Accumulation are optimal is because they respect maximal threshold-hitting times in the mean-preserving strategies. This allows us to emphasize that it is their optimality. The assumption that payoffs depend and not on which threshold is hit allows us to derive critical for the economic insights.

Indeed, we anticipate that many of the qualitative features of Pure Accumulation will persist under other examples, with other costs of learning, multiple state thresholds (there are already examples in the literature in Section 2). It is certainly possible to extend to accommodate these more general environments though. ¹²The advantage of our special setup is that it identifies strategies that are explicit and uniformly optimal.

¹¹Specifically, with a binary state and fixed thresholds, even the probability distribution over terminal beliefs by the market with multiple states. That is, all learning strategies yield the same constraint then ensures that all learning strategies that satisfy the same expected threshold-hitting times.

¹²Our solutions were based on a guess and verify approach that exploits the structure of our setup. In more general setups guessing the

moreover, allows us to isolate the role of time risk.

There are two promising avenues to explore in future work. First, how our results may extend to the case when the agent is time-risk averse. For these more general preferences, what is the optimal information acquisition? A second avenue is to extend our model of information acquisition into strategic settings with multiple agents in order to study the implications of flexible information acquisition.

References

- Arrow, K. J., Blackwell, D., & Girshick, M. A. (1953). Bayes' theorem and the theory of sequential decisions. *Journal of the American Statistical Association*, 48(262), 213–244.
- Caplin, A., Dean, M., & Leahy, J. (2017). Time discounting and time preference: Characterizing and Generalizing. *Working Paper 2365*, National Bureau of Economic Research.
- Che, Y.-K. & Miendorff, K. (2019). Optimal dynamic information acquisition. *American Economic Review*, 109(18), 2993–3029.
- Chen, M. K. (2013). The effect of language on economic savings rates, health behavior, and economic growth. *American Economic Review*, 103(2), 690–731.
- Chesson, H. W. & Viscusi, W. K. (2003). Commonalities in time discounting and time preference. *Journal of Human Capital*, 5(1), 57–71.
- Dasgupta, P. & Stiglitz, J. (1980). Uncertainty, industrial structure, and the nature of innovative activity. *The Bell Journal of Economics*, 1(2), 1–28.
- DeJarnette, P., Dillenberger, D., Gottlieb, D., & Smith, J. (2018). Time discounting and time preference: A review. *Journal of Economic Literature*, 56(2), 619–656.
- Ely, J., Frankel, A., & Kamienica, E. (2015). The role of time discounting in the theory of optimal information acquisition. *Journal of Economic Theory*, 151, 215–260.

- Frankel, A. & Kamenica, E. (2019). Quasi-rationality in American Economic Review 3650–80.
- Fudenberg, D., Strack, P., & Strzalecki, T. (2018). Timing of American Economic Review 121, 651–84.
- Georgiadis-Harris, A. (2023). Information acquisition.
- Gossner, O., Steiner, J., & Stewart, D. (2012). 89(4), 1717–1751.
- Hébert, B. & Woodford, M. (2017). Attention and sequential sampling. Technical report, National Bureau of Economic Research.
- Hébert, B. & Woodford, M. (2023). Rational inattention. Journal of Economic Theory 205, 612–647.
- Hörner, J. & Skrzypacz, A. (2017). Learning, experiments, and advances in Economics and Econometrics 13E-08.
- Karatzas, I. & Shreve, S. E. (1998). Brownian motion and stochastic calculus. Springer, 47–127.
- Ke, T. T. & Villas-Boas, J. M. (2019). Double tails in the Economic Theory 383–437.
- Lee, T. & Wilde, L. L. (1980). Market structure and The Quarterly Journal of Economics 94(2), 429–436.
- Liang, A., Mu, X., & Syrgkanis, V. (2019). Dynamic rationality. Xiv preprint arXiv:1910.07015.
- Loewenstein, G. & Prelec, D. (1992). Anomalies in time discounting and time preference. Quarterly Journal of Economics 107(2), 557–593.
- Matijka, F. & McKay, A. (2014). Rational inattention: a foundation for the multi-armed bandit problem. The American Economic Review 105(1), 272–298.
- Mayskaya, T. (2019). Dynamic choice of information. manuscript.

- Morris, S. & Strulik, W. (2019). *Time and the Equivalence of Sampling and Static Welfare: Implications for Costs*
- Moscarini, G. & Smith, L. (2001). *The Economics of Time-Dependent Risk*, *Journal of Economic Theory*, 69(6), 1629–1644.
- Müller, A. (1996). *Orderings of risks: A comparison of insurance*, *Mathematical Economics*, 17(3), 215–222.
- Nikandrova, A. & Pansy, R. (2018). *The Dynamics of Economic Growth*, *Journal of Economic Theory*, 133(1), 115–143.
- Onay, S. & Öncüler, A. (2007). *Intertemporal choice and the role of risk aversion*, *Journal of Economic Theory*, 134(2), 99–112.
- Protter, P. (2005). *Stochastic differential equations*
- Ratcliff, R. & Rouder, J. N. (1998). *Modeling response times*, *Psychological Review*, 105(1), 34–81.
- Reinganum, J. F. (1989). *The timing of innovation*, *Handbook of industrial organization*
- Sims, C. A. (2003). *Implications of a rational expectations model for macroeconomics*, *Journal of Economic Theory*, 50(3), 665–690.
- Steiner, J., Stewart, C., & Matijka, F. (2017). *Risk and delay in decision making*, *Journal of Economic Psychology*, 65(12), 521–553.
- Wald, A. (1947). *Foundations of a general theory of econometrics*, *Journal of the American Statistical Association*
- Zhong, W. (2022). *Optimal dynamic consumption and investment*, *Journal of Economic Theory*, 1537–1582.
- Zhong, W. & Bloedel, A. (2021). *The cost of optimal*